

Today's Topics

6.1 Gaussian & Laplacian pyramid construction

6.2 Applications: image blending,
editing, texture synthesis

Topic 6:

Hierarchical image representations

1. Gaussian & Laplacian pyramids
2. Applications:
 1. Multi-resolution image blending
 2. Multi-resolution image editing
 3. Multi-resolution texture synthesis

Topic 6.1:

Gaussian & Laplacian Pyramids

- The Gaussian pyramid (intro)
- The convolution operation
- Constructing the gaussian pyramid
 - The REDUCE() function
- Constructing the Laplacian pyramid
 - The EXPAND() function

The Gaussian Pyramid

The Gaussian Pyramid: A representation in multiple scales



Original Image

The Gaussian Pyramid

The goal is to define a representation in which image information at different scales is explicitly available (i.e. does not need to be computed when needed)

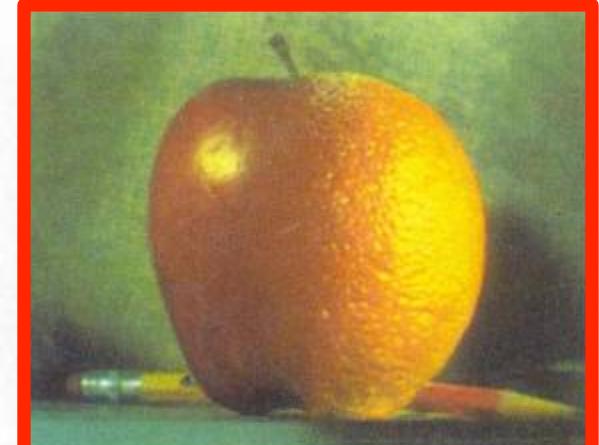
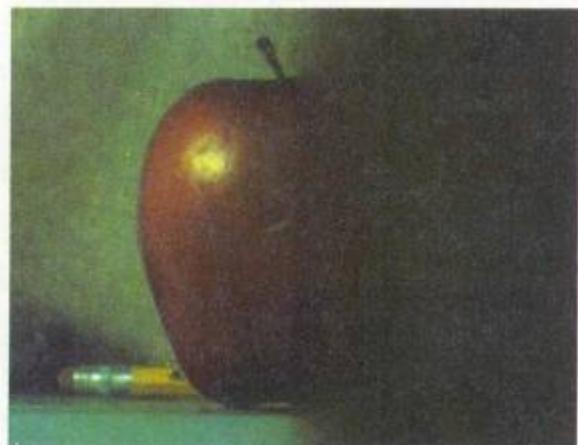
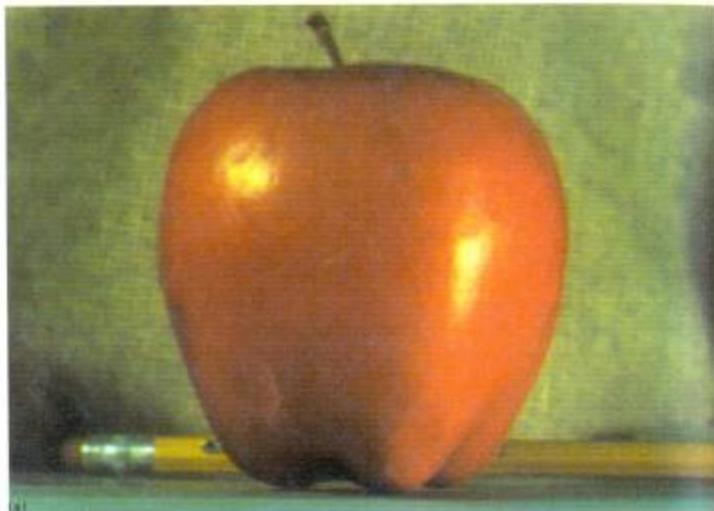
Applications:

- Scale invariant template matching (like faces)
- Progressive image transmission
- Image blending
- Efficient feature search
- ...

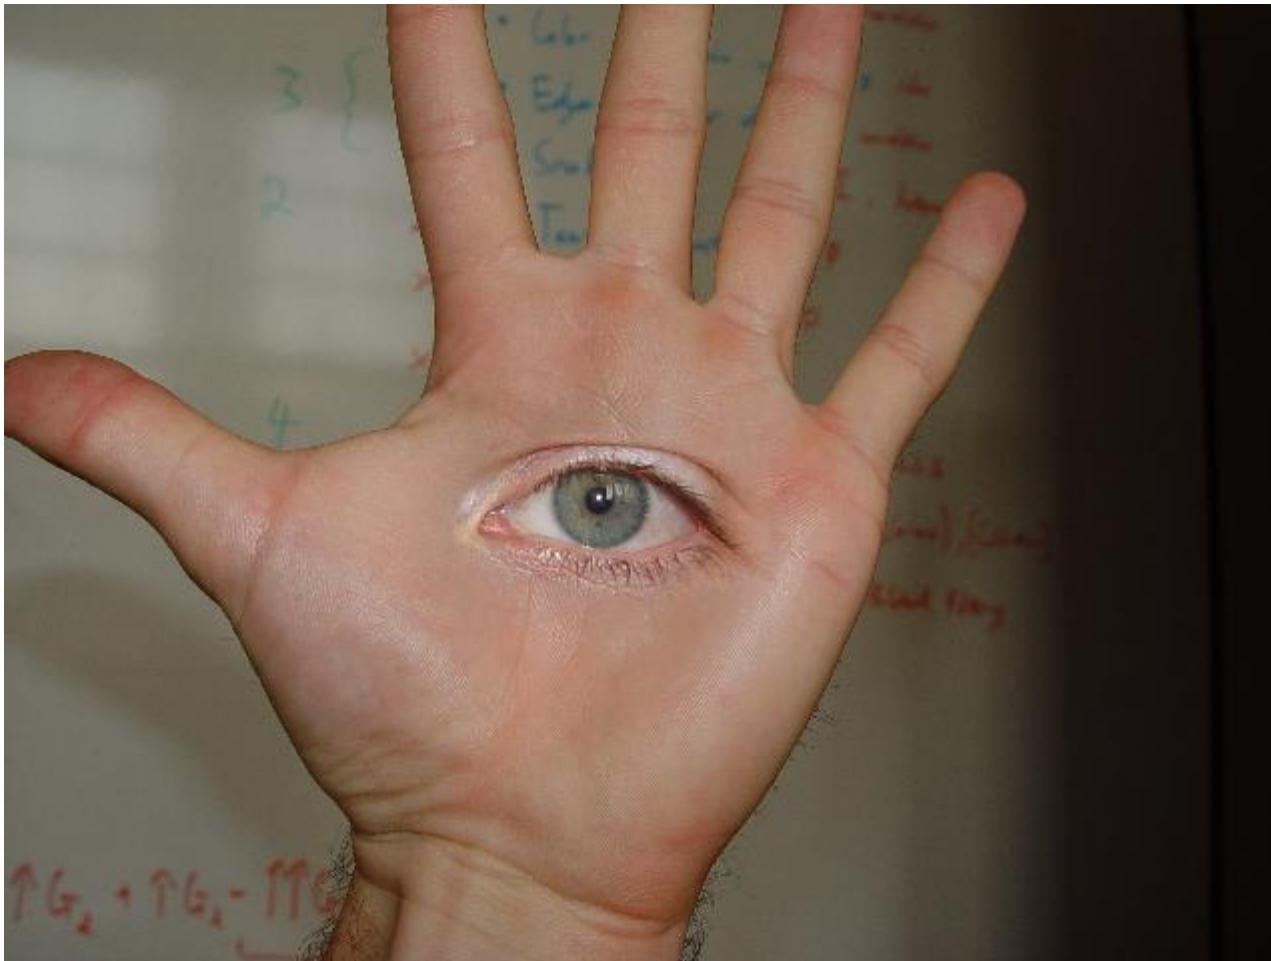


Application 1: Pyramid Image Blending

Goal: Merge two images without visible seams



Horror Photo

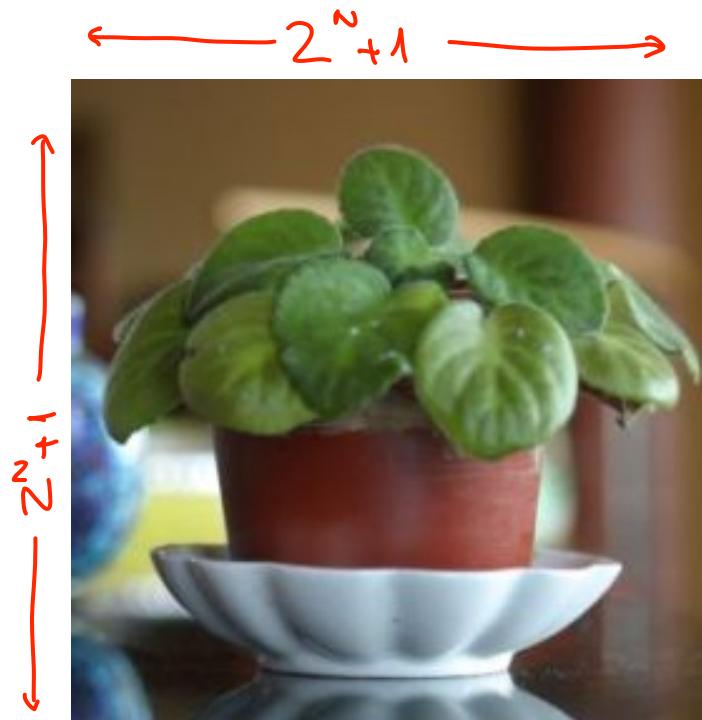


© prof. dmartin

The Gaussian Pyramid

The elements of a Gaussian Pyramids are smoothed copies of the image at different scales.

Input: Image I of size $(2^N+1) \times (2^N+1)$



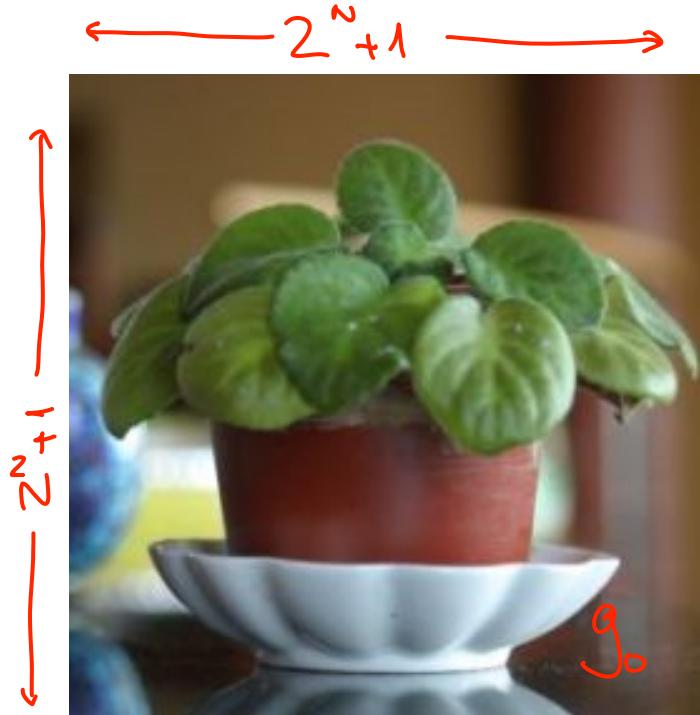
The Gaussian Pyramid

The elements of a Gaussian Pyramids are smoothed copies of the image at different scales.

Input: Image I of size $(2^N+1) \times (2^N+1)$

Output: Images g_0

Note: The original image is part of the output!



The Gaussian Pyramid

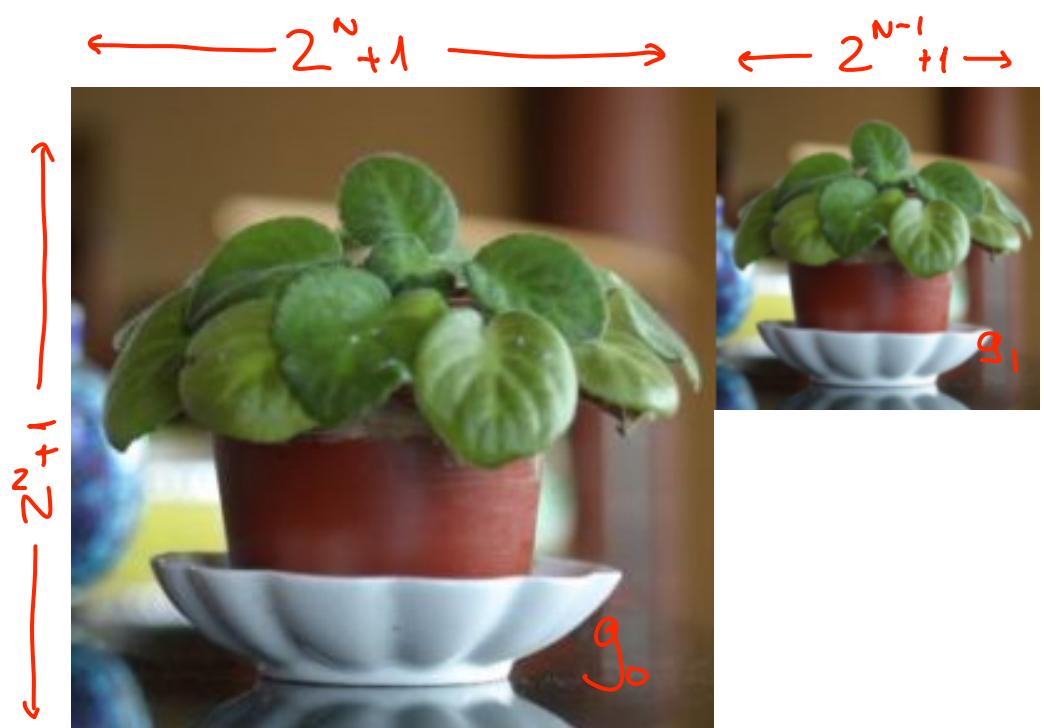
The elements of a Gaussian Pyramids are smoothed copies of the image at different scales.

Input: Image I of size $(2^N+1) \times (2^N+1)$

Output: Images g_0, g_1

where the size of g_1 is:

$$(2^{N-1}+1) \times (2^{N-1}+1)$$



The Gaussian Pyramid

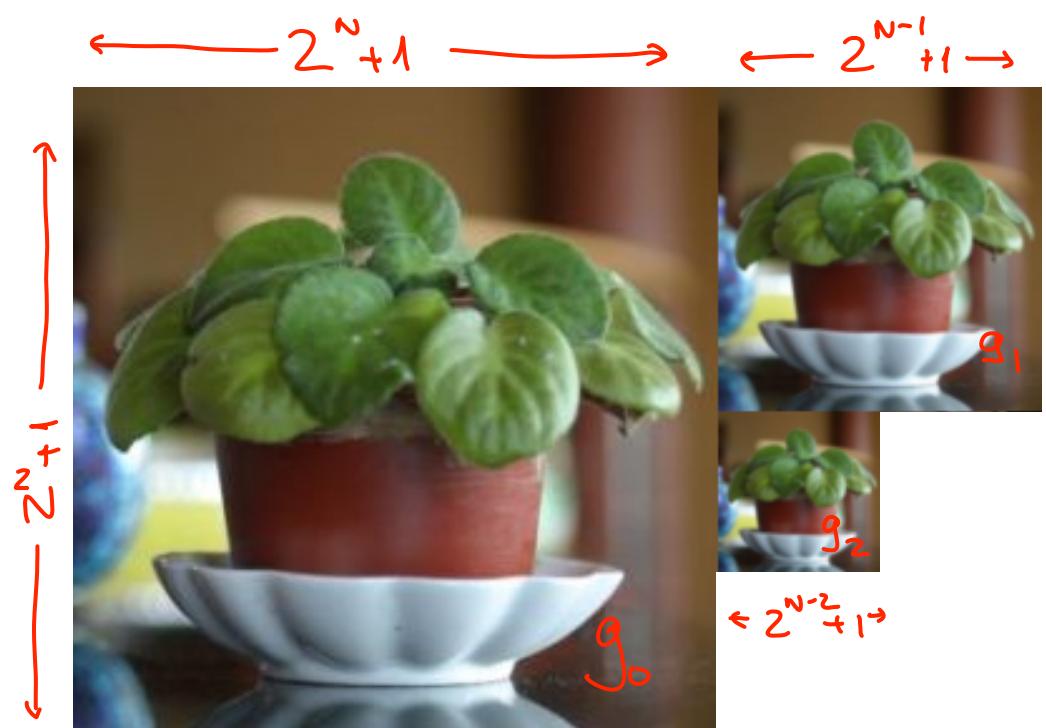
The elements of a Gaussian Pyramids are smoothed copies of the image at different scales.

Input: Image I of size $(2^N+1) \times (2^N+1)$

Output: Images g_0, g_1, g_2

where the size of g_2 is:

$$(2^{N-2}+1) \times (2^{N-2}+1)$$



The Gaussian Pyramid

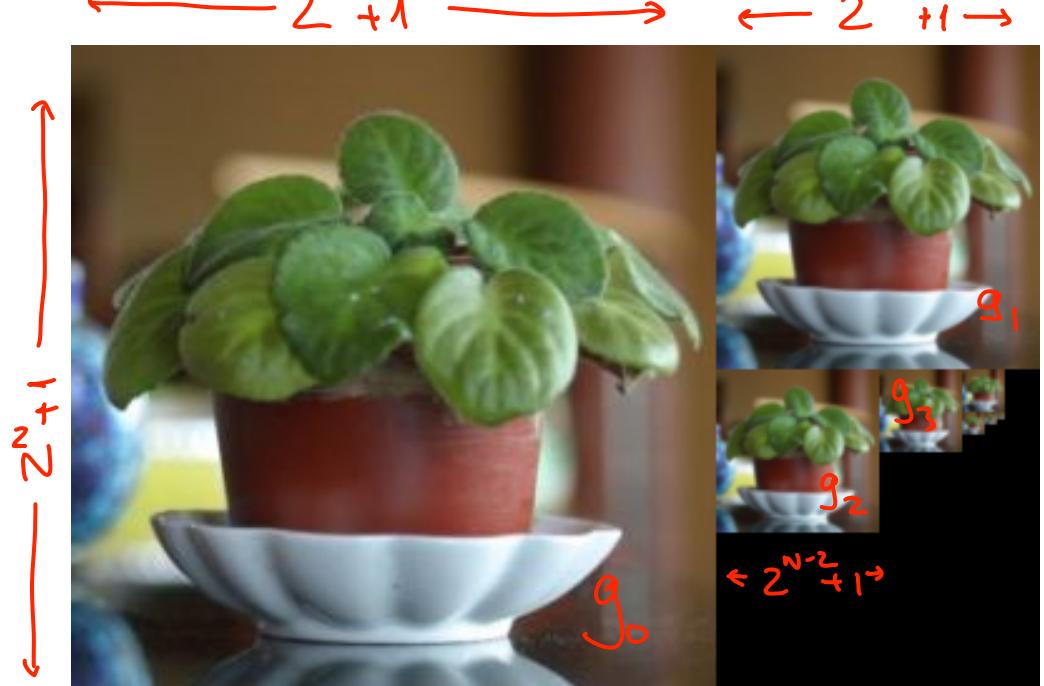
The elements of a Gaussian Pyramids are smoothed copies of the image at different scales.

Input: Image I of size $(2^N+1) \times (2^N+1)$

Output: Images g_0, g_1, \dots, g_{N-1}

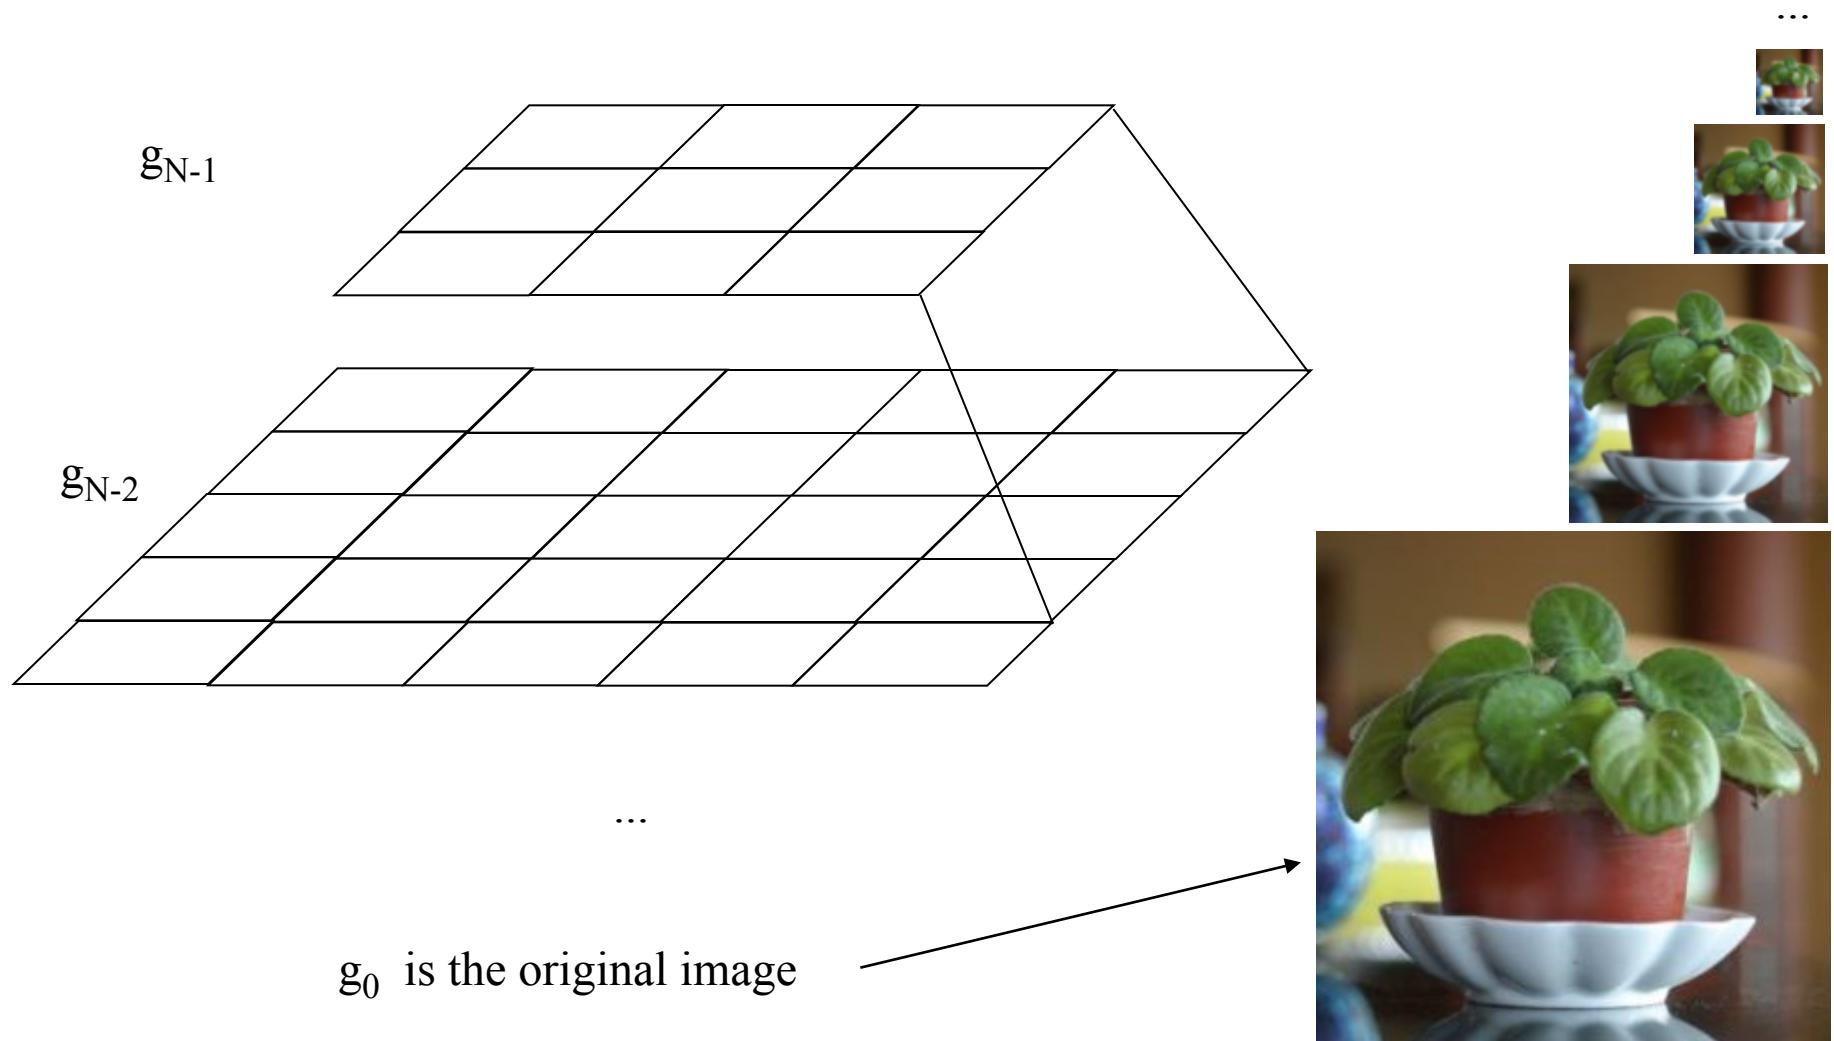
where the size of g_i is:

$$(2^{N-i}+1) \times (2^{N-i}+1)$$

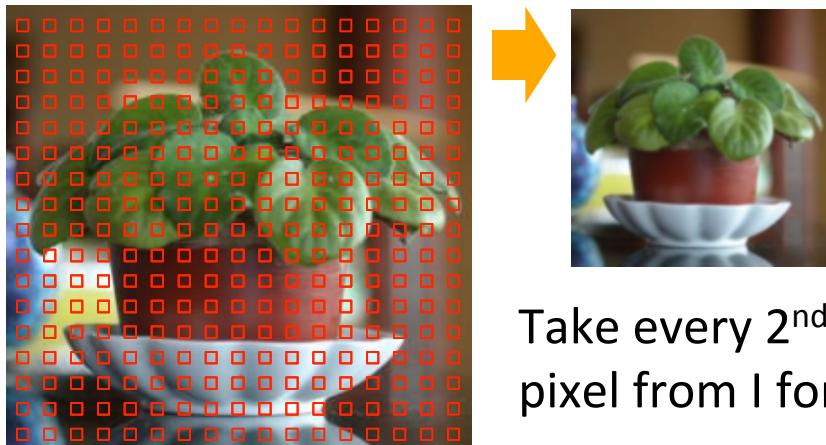


And they called this a Pyramid?

Yes, because the representation can be pictured as a pyramid of $3 \times 3, 5 \times 5, 9 \times 9, \dots, (2^N+1) \times (2^N+1)$ images when stacked.

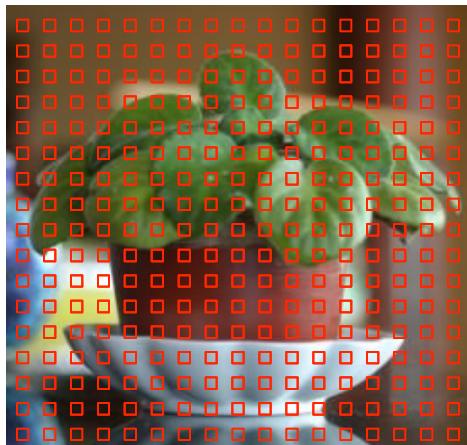


This is silly, why not just sub-sample?

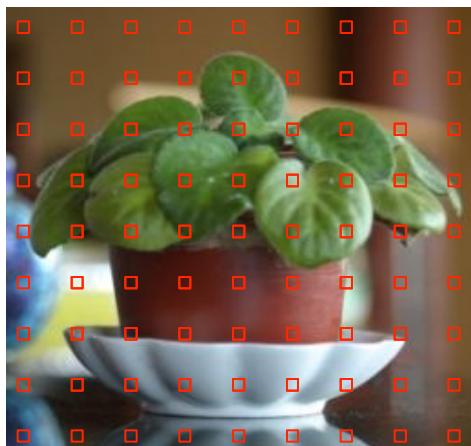


Take every 2nd
pixel from I for g_1

This is silly, why not just sub-sample?

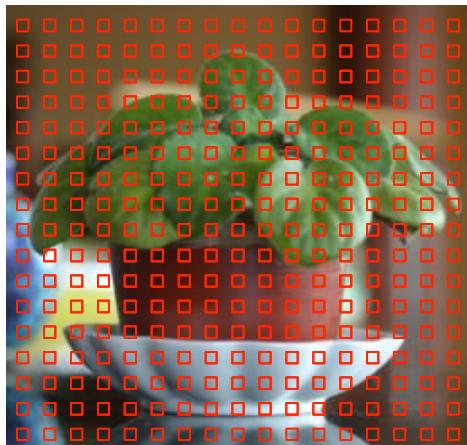


Take every 2nd
pixel from I for g_1

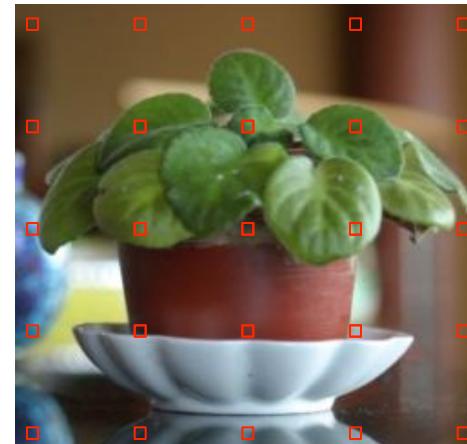


Take every 4th
pixel from I for g_2

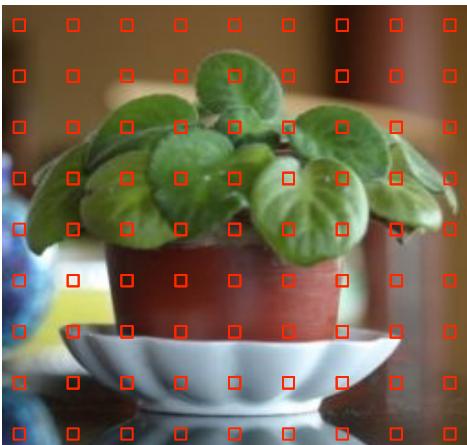
This is silly, why not just sub-sample?



Take every 2nd pixel from I for g_1



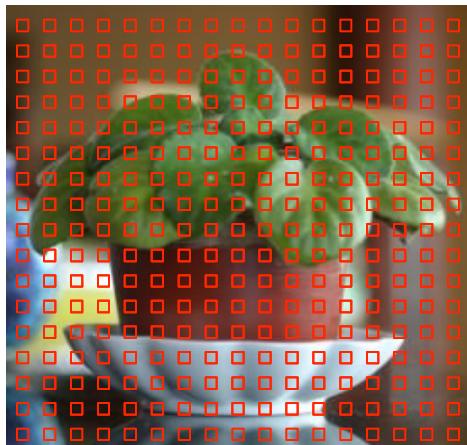
Take every 8th pixel from I for g_3



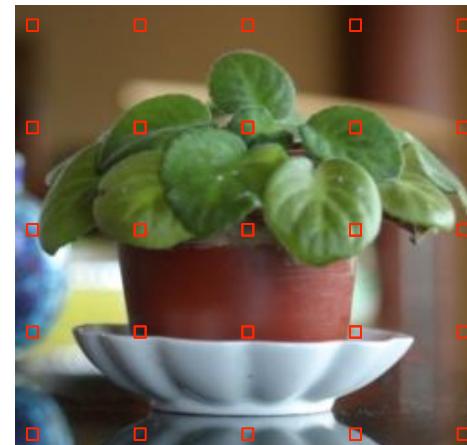
Take every 4th pixel from I for g_2

Final result

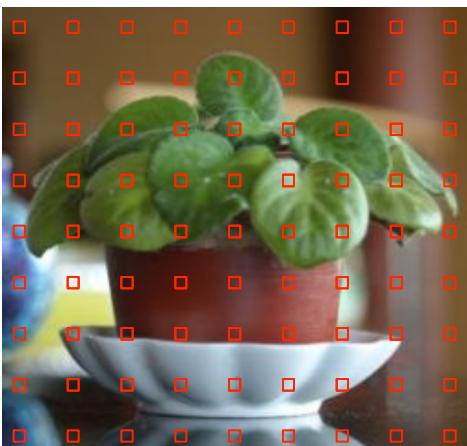
This is silly, why not just sub-sample?



Take every 2nd pixel from I c



Take every 8th pixel from I for g_3



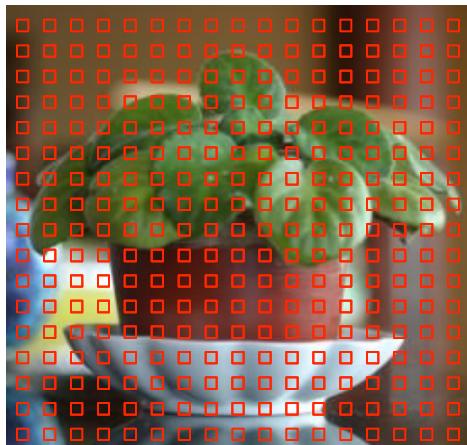
Take every 4th pixel from I for g_2



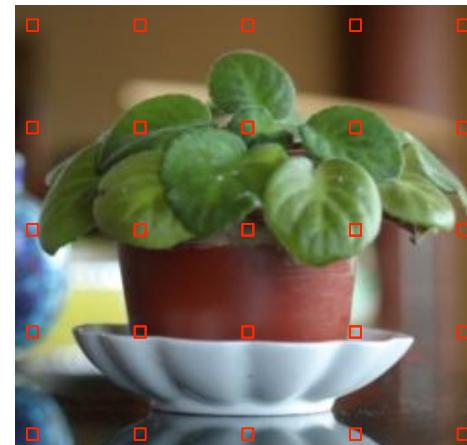
Take every 16th pixel from I for g_4

and so on...

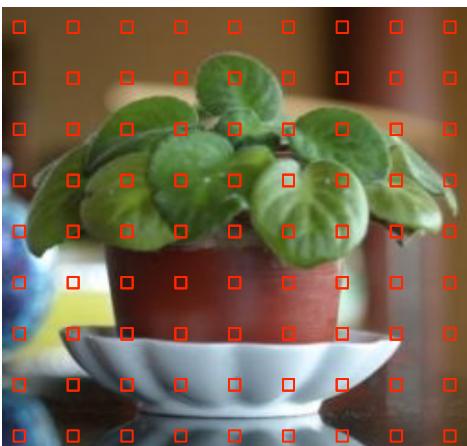
This is silly, why not just sub-sample?



Take every 2nd pixel from I for g_1



Take every 8th pixel from I for g_3



Take every 4th pixel from I for g_2



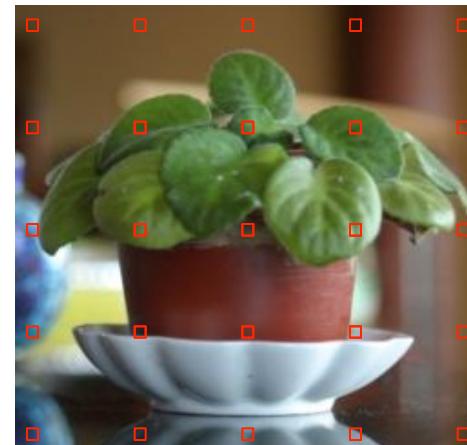
Take every 16th pixel from I for g_4

All the information is there already, or is it not?

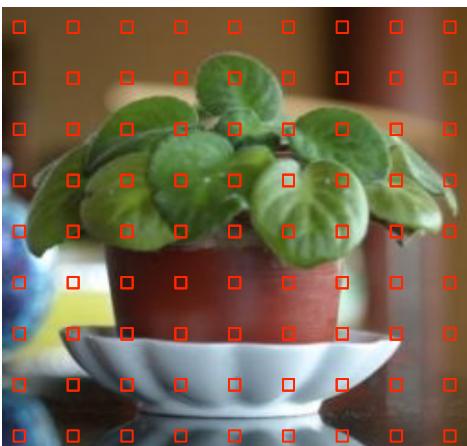
This is silly, why not just sub-sample?



Take every 2nd pixel from I for g_1



Take every 8th pixel from I for g_3



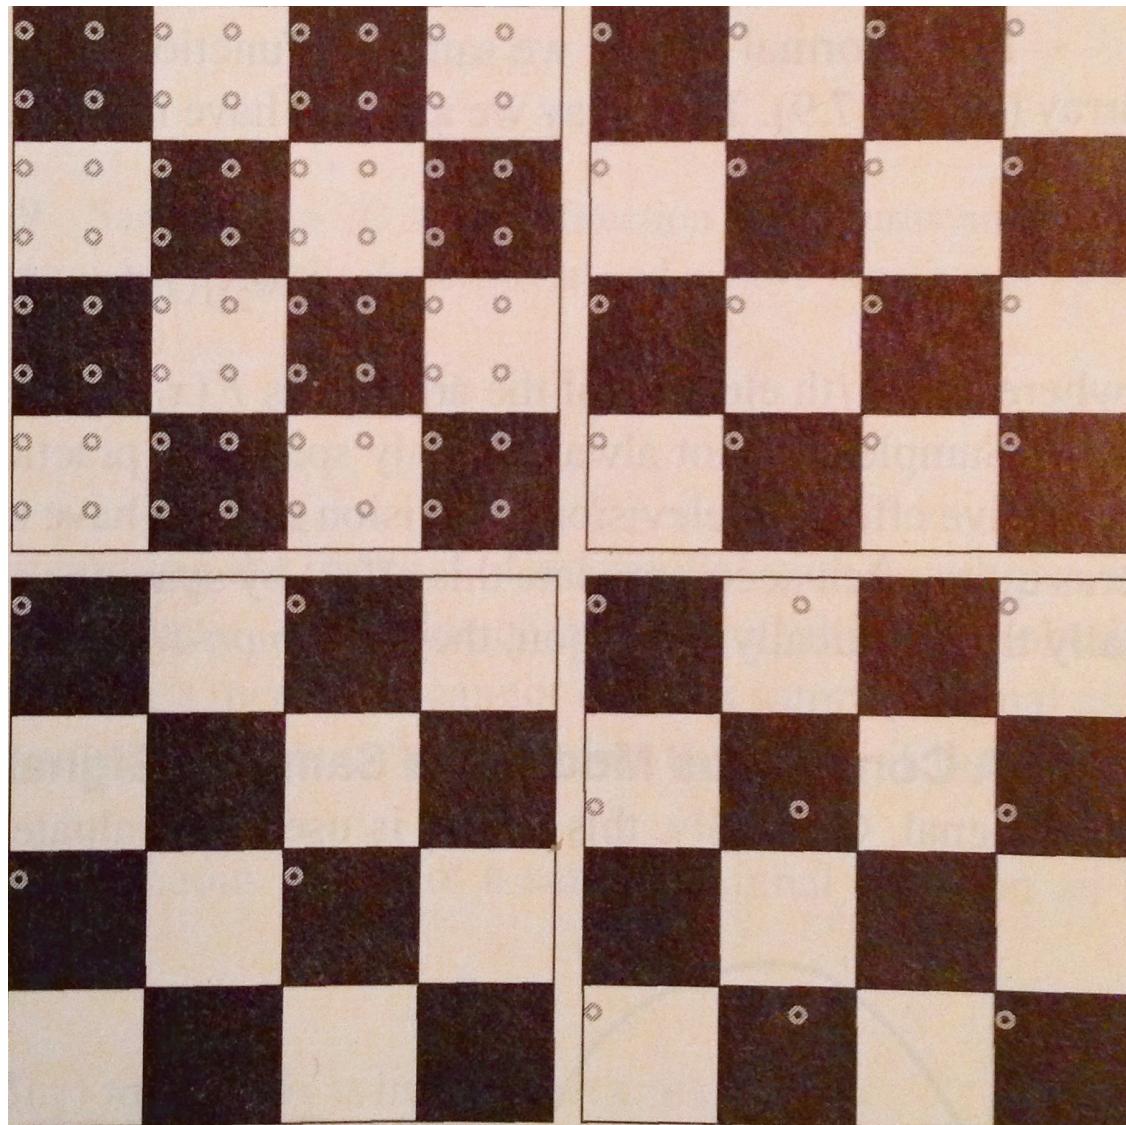
Take every 4th pixel from I for g_2



Take every 16th pixel from I for g_4

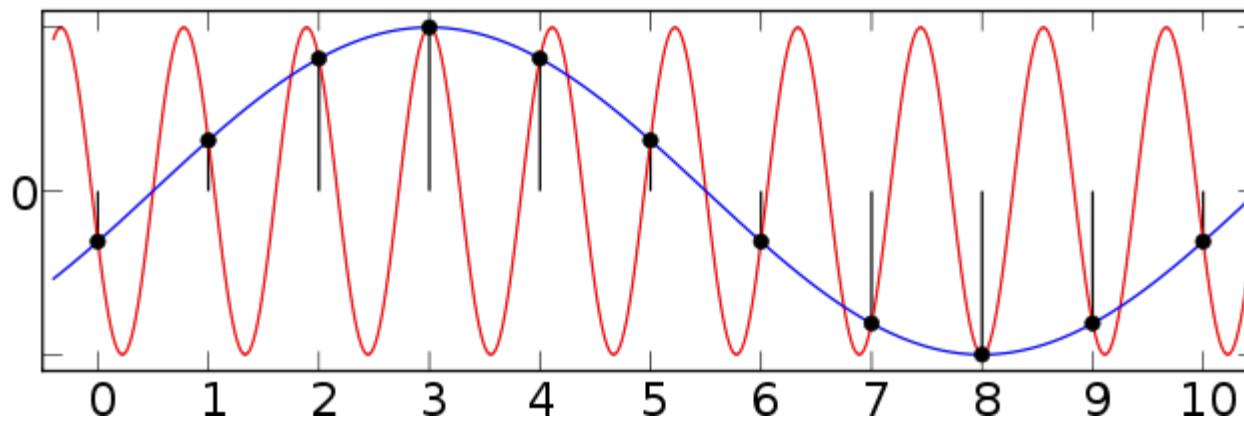
NO! SUBSAMPLING ALONE LEADS TO ALIASING

Aliasing?



Forsyth and Ponce

Sub-sampling is not enough!

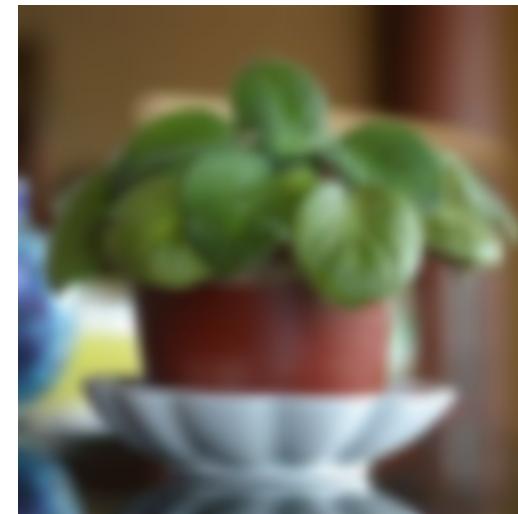


Because of “Aliasing”, sub-sampling may have catastrophic effects.

Aliasing arises when a signal is sampled at a rate that is insufficient to capture the changes in the signal (and in Gaussian Pyramids, this will always happen as sub-sampling is twice as sparse at each level!)

Smoothing

The solution to aliasing effects is smoothing, effectively reducing the maximum frequency of image features. In other words smoothing removes the fast changes that sub-sampling would miss.



Smoothing

Building (or computing) a Gaussian Pyramid

To generate a Gaussian pyramid,
iterate between these two steps:

Smoothing: Remove
high-frequency
components that
could cause aliasing.

Down-sampling:
Reduce the image size
by $\frac{1}{2}$ at each level.



Building (or computing) a Gaussian Pyramid

To generate a Gaussian pyramid,
iterate between these two steps:

Smoothing: Remove
high-frequency
components that
could cause aliasing.

Down-sampling:
Reduce the image size
by $\frac{1}{2}$ at each level.



Good, but how do we actually implement this?

Topic 6.1:

Gaussian & Laplacian Pyramids

- The Gaussian pyramid (intro)
- The convolution operation
- Constructing the gaussian pyramid
 - The REDUCE() function
- Constructing the Laplacian pyramid
 - The EXPAND() function

Image Smoothing Using Averaging Masks

We know about (both the normalized and un-normalized) the cross correlation operators.

$$CC(x_i, T) = X_i^T \cdot T$$

$$NCC(x_i, T) = \frac{X_i^T \cdot T}{\|x_i\| \cdot \|T\|}$$

What happens if we evaluate them with

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

or

1/25	1/25	1/25	1/25	1/25
1/25	1/25	1/25	1/25	1/25
1/25	1/25	1/25	1/25	1/25
1/25	1/25	1/25	1/25	1/25
1/25	1/25	1/25	1/25	1/25

everywhere in the image?

Image Smoothing Using Averaging Masks

Smoothing can be achieved by averaging neighboring pixels.

The strength of a smoothing operator is proportional to the number of pixels it averages.

Averaging can be computed as the Cross-Correlation of the image with a constant kernel, like:

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

or

1/25	1/25	1/25	1/25	1/25
1/25	1/25	1/25	1/25	1/25
1/25	1/25	1/25	1/25	1/25
1/25	1/25	1/25	1/25	1/25
1/25	1/25	1/25	1/25	1/25

Image Smoothing Using Averaging Masks

Original Image



Image Smoothing Using Averaging Masks

Result of Cross-Correlation with 3x3 Mask



Image Smoothing Using Averaging Masks

Result of Cross-Correlation with 5x5 Mask



1/25	1/25	1/25	1/25	1/25
1/25	1/25	1/25	1/25	1/25
1/25	1/25	1/25	1/25	1/25
1/25	1/25	1/25	1/25	1/25
1/25	1/25	1/25	1/25	1/25

Image Smoothing Using Averaging Masks

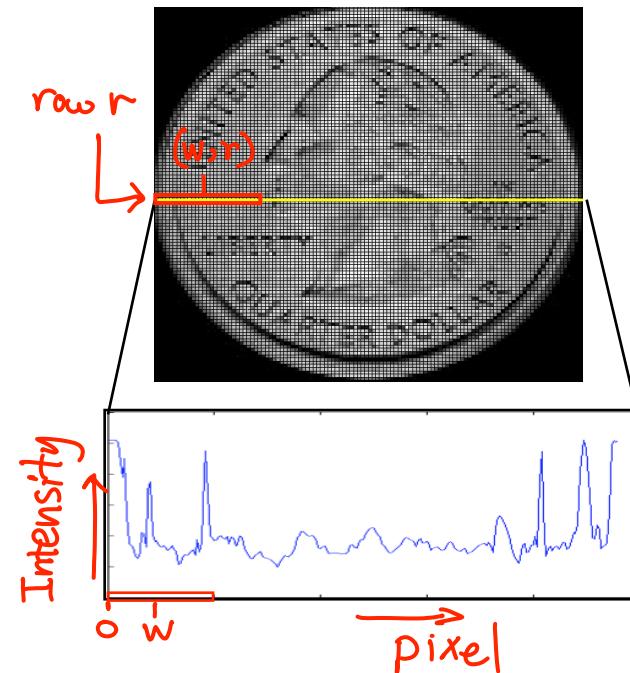
Result of Cross-Correlation with 15x15 Mask



Image Cross Correlation \Leftrightarrow Matrix Multiplication

Cross correlation in 1D can be computed using matrix multiplication, for instance, let:

I: one row of the image (with M pixels)



and

T: a template (with $2w+1$ pixels), such as $[T_{-w}, T_{-w-1}, \dots, T_0, \dots, T_{w-1}, T_w]$

Image Cross Correlation \Leftrightarrow Matrix Multiplication

Then, the cross correlation at pixel w can be computed using:

I: one row of the image (with M pixels)

T: a template (with $2w+1$ pixels), such as $[T_{-w}, T_{-w-1}, \dots, T_0, \dots, T_{w-1}, T_w]$

$$CC(I, T) = [T_{-w} \ T_{-w+1} \ \dots \ T_0 \ \dots \ T_w \ 0 \ \dots \ 0]$$

$$\begin{bmatrix} I_0 \\ I_1 \\ I_2 \\ \vdots \\ I_w \\ \vdots \\ I_{2w} \\ I_{2w+1} \\ \vdots \\ I_{M-1} \end{bmatrix}$$

Image Cross Correlation \Leftrightarrow Matrix Multiplication

Then, the cross correlation at pixel w can be computed using:

I: one row of the image (with M pixels)

T: a template (with $2w+1$ pixels), such as $[T_{-w}, T_{-w-1}, \dots, T_0, \dots, T_{w-1}, T_w]$

And then the next pixel:

pixel w

$$\hookrightarrow [T_{-w} \ T_{-w+1} \ \dots \ T_0 \ \dots \ T_w \ 0 \ \dots \ 0]$$

$$\mapsto [0 \ T_{-w} \ T_{-w+1} \ \dots \ T_{w-1} \ T_w \ \dots \ 0]$$

pixel $w+1$

$$\begin{bmatrix} I_0 \\ I_1 \\ I_2 \\ \vdots \\ I_w \\ \vdots \\ I_{2w} \\ I_{2w+1} \\ \vdots \\ I_{M-1} \end{bmatrix}$$

Image Cross Correlation \Leftrightarrow Matrix Multiplication

Then, the cross correlation at pixel w can be computed using:

I : one row of the image (with M pixels)

T : a template (with $2w+1$ pixels), such as $[T_{-w}, T_{-w-1}, \dots, T_0, \dots, T_{w-1}, T_w]$

And for all the pixels in $[w, M-w-1]$:

The diagram shows a matrix multiplication between a column vector I and a row vector T . The column vector I has entries labeled $I_0, I_1, I_2, \dots, I_w, \dots, I_{2w}, I_{2w+1}, \dots, I_{M-1}$. The row vector T has entries labeled $T_{-w}, T_{-w+1}, \dots, T_0, \dots, T_w, 0, \dots, 0$. Dashed arrows point from each entry in I to its corresponding position in the product matrix, showing that each row is a 1-pixel right shift of the row above it. The resulting matrix is a **TOEPLITZ matrix**.

each row is a 1-pixel right shift of the row above it
(called a TOEPLITZ matrix)

Image Cross Correlation \Leftrightarrow Matrix Multiplication

Then, the cross correlation at pixel w can be computed using:

I : one row of the image (with M pixels)

T : a template (with $2w+1$ pixels), such as $[T_{-w}, T_{-w-1}, \dots, T_0, \dots, T_{w-1}, T_w]$

And for all the pixels in $[w, M-w-1]$:

A diagram illustrating a matrix multiplication for image cross-correlation. On the left, a vertical vector I is shown with components $I_0, I_1, I_2, \dots, I_w, \vdots, I_{2w}, I_{2w+1}, \vdots, I_{M-1}$. To its right is a large square matrix. The top row of the matrix has entries $T_{-w}, T_{-w+1}, \dots, T_0, \dots, T_w, 0, \dots, 0$. Dashed arrows point from each entry I_i in the vector I to the corresponding position in the first row of the matrix. Below the matrix, handwritten text states: "each row is a 1-pixel right shift of the row above it (called a TOEPLITZ matrix)".

How about pixels in:

[1, 2, ..., w]

and

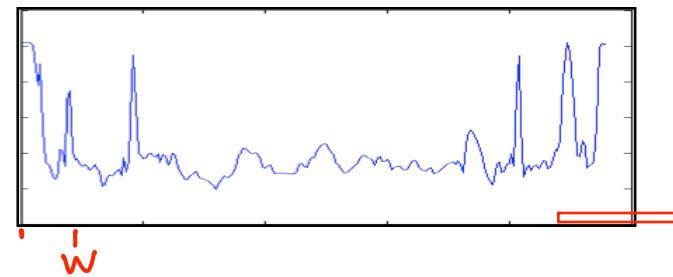
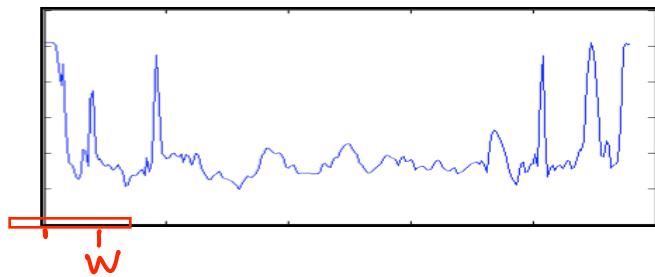
[$M-w, M-w+1, \dots, M$]?

Image Cross Correlation \Leftrightarrow Matrix Multiplication

How about pixels at $[1, 2, \dots, w]$ and $[M-w, M-w+1, \dots, M]$?

Options (to taste, with advantages and disadvantages):

- Wrap around in I
- Define new templates
- Assume I is zero for the out-of-image
- Extrapolate (in-paint!?)



Cross-Correlation Expressed as a Sum

$$J = CC(I, T) = \begin{bmatrix} J_0 \\ J_1 \\ \vdots \\ J_w \\ J_{w+1} \\ \vdots \\ J_{M-1} \end{bmatrix} = \begin{bmatrix} T_0 T_1 \dots T_w 0 \dots 0 \\ T_1 T_0 T_1 \dots T_w 0 \\ \vdots \\ T_{-w} T_0 \dots T_w 0 \dots 0 \\ 0 T_{-w} \dots T_0 \dots T_{w-1} T_w \dots 0 \\ \vdots \\ 0 0 T_{-w} T_{w-1} \dots T_{-1} T_0 T_1 \\ 0 \dots 0 T_{-w} \dots T_2 T_{-1} T_0 \end{bmatrix} \begin{bmatrix} I_0 \\ I_1 \\ \vdots \\ I_w \\ I_{w+1} \\ \vdots \\ I_{2w} \\ \vdots \\ I_{M-1} \end{bmatrix}$$

Cross-Correlation Expressed as a Sum

$$J = CC(I, T) = \begin{bmatrix} J_0 \\ J_1 \\ \vdots \\ J_w \\ J_{w+1} \\ \vdots \\ J_{M-1} \end{bmatrix} = \begin{bmatrix} T_0 & T_1 & \cdots & T_w & 0 & \cdots & 0 \\ T_{-1} & T_0 & T_1 & \cdots & T_w & 0 & \\ T_{-w} & T_{-w+1} & \cdots & T_0 & -T_w & 0 & 0 \\ 0 & T_{-w} & T_0 & \cdots & -T_{w-1} & T_w & \cdots & 0 \\ 0 & 0 & T_w & T_{w-1} & \cdots & T_{-1} & T_0 & T_1 \\ 0 & \cdots & 0 & T_{-w} & \cdots & T_2 & T_1 & T_0 \end{bmatrix} \begin{bmatrix} I_0 \\ I_1 \\ \vdots \\ I_w \\ I_{w+1} \\ \vdots \\ I_{M-1} \end{bmatrix}$$

General
sum notation:

$$J_i = \sum_{k=0}^{M-1} I_k T_{k-i}$$

$0 \leq i \leq M-1$

Cross-Correlation Expressed as a Sum

$$J = CC(I, T) = \begin{bmatrix} J_0 \\ J_1 \\ \vdots \\ J_w \\ J_{w+1} \\ \vdots \\ J_{M-1} \end{bmatrix} = \begin{bmatrix} T_0 & T_1 & \cdots & T_w & 0 & \cdots & 0 \\ T_1 & T_0 & T_1 & \cdots & T_w & 0 & \\ T_w & \cdots & T_0 & \cdots & T_w & 0 & \\ 0 & T_{-w} & T_0 & \cdots & T_{w-1} & T_w & \cdots & 0 \\ 0 & 0 & T_w & T_{w-1} & \cdots & T_1 & T_0 & \\ 0 & \cdots & 0 & T_{-w} & \cdots & T_2 & T_1 & T_0 \end{bmatrix} \begin{bmatrix} I_0 \\ I_1 \\ \vdots \\ I_w \\ I_{w+1} \\ \vdots \\ I_{M-1} \end{bmatrix}$$

General sum notation:

$$J_i = \sum_{k=0}^{M-1} I_k T_{k-i}$$

$0 \leq i \leq M-1$

for instance:

$$J_w = \sum_{k=0}^{M-1} I_k \cdot T_{k-w}$$

$$J_{w+1} = \sum_{k=0}^{M-1} I_k \cdot T_{k-w-1}$$

$$= \sum_{k=0}^{M-1} I_k \cdot T_{k-(w+1)}$$

Indexing by image position, not template position

$$J = CC(I, T) = \begin{bmatrix} J_0 \\ J_1 \\ \vdots \\ J_w \\ J_{w+1} \\ \vdots \\ J_{M-1} \end{bmatrix} = \begin{bmatrix} T_0 T_1 - T_w & 0 & \cdots & 0 \\ T_1 T_0 T_1 - T_w & 0 & & \\ \vdots & & & \\ T_w T_0 - T_w & 0 & \cdots & 0 \\ 0 & T_w & T_0 - T_w & T_w & \cdots & 0 \\ & & & & & \\ 0 & 0 & T_w T_{w+1} - T_1 T_0 & T_1 T_0 & \cdots & \\ 0 & \cdots & 0 & T_w - T_2 T_1 T_0 & & \end{bmatrix} \begin{bmatrix} I_0 \\ I_1 \\ \vdots \\ I_w \\ I_{w+1} \\ \vdots \\ I_{M-1} \end{bmatrix}$$

General sum notation:

$$J_i = \sum_{k=0}^{M-1} I_k T_{k-i}$$

$0 \leq i \leq M-1$

Equivalent to:

$$J_i = \sum_{l=0}^{M-1} I_{l+i} T_l$$

$0 \leq i \leq M-1$

* obtained by substituting
 $l = k-i$ in General Sum
 Notation formula

The Convolution Operation

A similar operation to Cross Correlation is Convolution:

Cross-correlation:

$$J_i = \sum_{k=0}^{M-1} I_k \cdot T_{k-i}$$

Convolution:

$$(I * T)_i = \sum_{k=0}^{M-1} I_k \cdot T_{i-k}$$

$$[T_{-w} \quad T_{-w+1} \dots T_0 \dots T_w \quad 0 \dots 0]$$

$$[T_w \quad T_{w-1} \dots T_0 \dots T_{-w} \quad 0 \dots 0]$$

$$\begin{bmatrix} I_0 \\ I_1 \\ I_2 \\ \vdots \\ I_w \\ \vdots \\ I_{2w} \\ I_{2w+1} \\ \vdots \\ I_{M-1} \end{bmatrix}$$

The Convolution Operation

The convolution operation is one of the most fundamental operations in image (and signal) processing.

In this context

I is the “image” or the “signal”

T is the “filter”, “mask”, “template”, “impulse response”,
“kernel”...

Notation: $I * T$: reads as “the convolution of I with T ” and means:

Convolution:

$$(I * T)_i = \sum_{k=0}^{m-1} I_k \cdot T_{i-k}$$

The Convolution Properties

1. For symmetric masks, convolution is equal to cross-correlation:

$$CC(I, T) = I * T$$

when $T_i = T_{-i}$

2. Commutativity:

$$I * T = T * I$$

3. Linearity:

$$(aI + bJ) * T =$$
$$a(I * T) + b(J * T)$$

for any constants a, b

You may want
to prove
these as an
exercise.

The Convolution Properties

But images are 2D!

The Convolution Properties

Similar to 2D cross-correlation. For a M by N image:

$$(I * T)_{(i,j)} = \sum_{k=0}^{M-1} \sum_{l=0}^{N-1} I(k,l) \cdot T(i-k, j-l)$$

This can get expensive. If the image is M by N and the Template is P by Q, the complexity is:

$$O(MNPQ)$$

The Separable Convolution

There is one very special case, when the Template is “separable”.

Separable templates are such that $T = PQ^T$ for some vectors P, Q .

$$\begin{array}{c} \text{equal to} \\ P_i \cdot Q_j \end{array} \quad \left[\begin{array}{c} \text{i-th row} \\ \cdots \\ \cdots T(i,j) \\ \cdots \\ \text{j-th col} \end{array} \right] \approx \left[\begin{array}{c} P_i \\ \vdots \\ \vdots \end{array} \right] \left[\begin{array}{c} \text{j-th col} \\ \vdots \\ Q_j \\ \vdots \\ \text{vector } Q \end{array} \right]$$

vector P

The Separable Convolution

There is one very special case, when the Template is “**separable**”.

For instance, is the “Sobel” kernel separable?

$$\text{Sobel} = \begin{pmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{pmatrix} = \begin{pmatrix} & & \end{pmatrix} \begin{pmatrix} & & \end{pmatrix}$$

The Separable Convolution

There is one very special case, when the Template is “**separable**”.

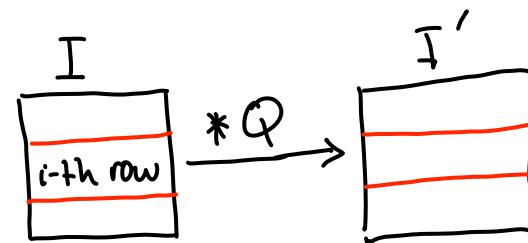
For instance, is the “Sobel” kernel separable?

$$\text{Sobel} = \begin{pmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{pmatrix} = \begin{pmatrix} 1 \\ 2 \\ 1 \end{pmatrix} \begin{pmatrix} [-1 \ 0 \ 1] \end{pmatrix}$$

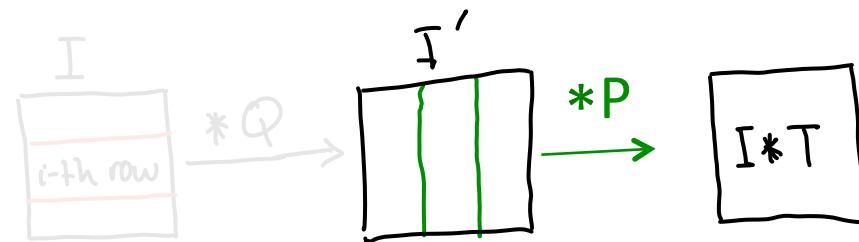
The Convolution Operation

When a kernel is separable the 2D convolution can be obtained from 2, cascaded 1D convolutions. The algorithm is as follows:

1. Compute the 1D convolution between each row of I and Q to obtain I' .



2. Compute the 1D convolution between each row of I' and P , to obtain the final 2D result:



The Convolution Operation

When a kernel is separable the 2D convolution can be obtained from 2, cascaded 1D convolutions.

In the previous slide, we first convolved rows of I with Q to obtain I' and then columns of I' with P to obtain the final result.

The same result can be obtained by first convolving columns of I with P to obtain I'' and then convolve rows of I'' with Q .

You may want to prove this as an exercise

Topic 6.1:

Gaussian & Laplacian Pyramids

- The gaussian pyramid (intro)
- The convolution operation
- **Constructing the gaussian pyramid**
 - The REDUCE() function
- Constructing the Laplacian pyramid
 - The EXPAND() function

The Gaussian Pyramid

The Gaussian Pyramid: A representation in multiple scales



Original Image

The Gaussian Pyramid

The elements of a Gaussian Pyramids are smoothed copies of the image at different scales.

Input: Image I of size $(2^N+1) \times (2^N+1)$

Output: Images g_0, g_1, \dots, g_{N-1}

where the size of g_i is:

$$(2^{N-i}+1) \times (2^{N-i}+1)$$



Building (or computing) a Gaussian Pyramid

To generate a Gaussian pyramid,
iterate between these two steps:

Smoothing: Remove
high-frequency
components that
could cause aliasing.

Down-sampling:
Reduce the image size
by $\frac{1}{2}$ at each level.



Operation #1: Smooth Image, recursively

Take the original image g_0 and compute g_1 using:

$$\hat{g}_1(i,j) = \sum_{m=-2}^2 \sum_{n=-2}^2 w(m,n) \cdot g_0(i-m, j-n)$$



$\hat{g}_1 = w * g_0$

a 5×5 filter

like

1/25	1/25	1/25	1/25	1/25
1/25	1/25	1/25	1/25	1/25
1/25	1/25	1/25	1/25	1/25
1/25	1/25	1/25	1/25	1/25
1/25	1/25	1/25	1/25	1/25

Operation #1: Smooth Image at N-1 Scales

To estimate \hat{g}_2 repeat using \hat{g}_1 as the input:

$$\hat{g}_2 = w * \hat{g}_1$$

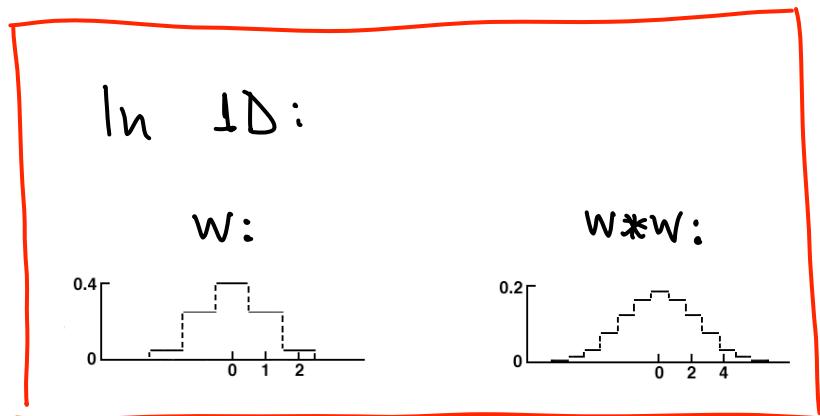
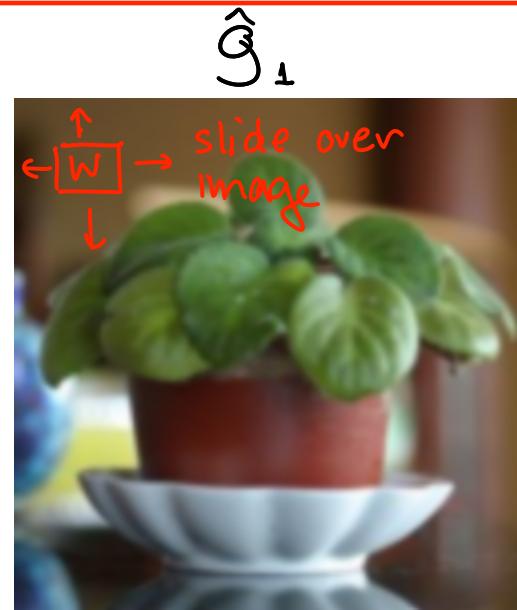
$$= w * (w * g_0)$$

$$= \underbrace{(w * w)}_{\text{can be thought of as a filter}} * g_0$$

can be thought of as a filter

$$h = w * w$$

whose radius is twice that of w

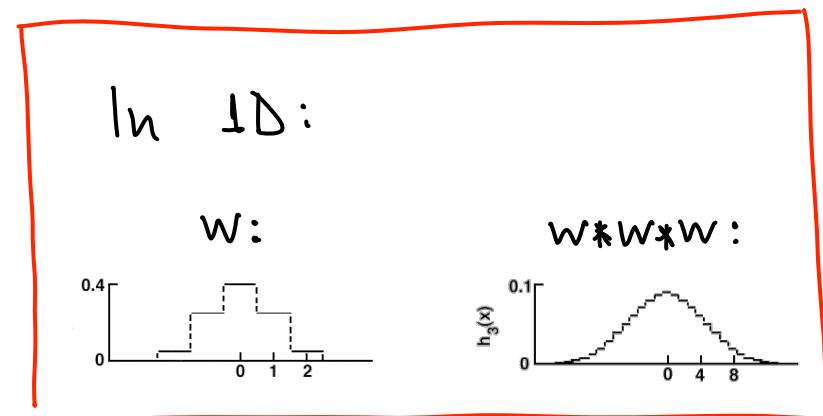
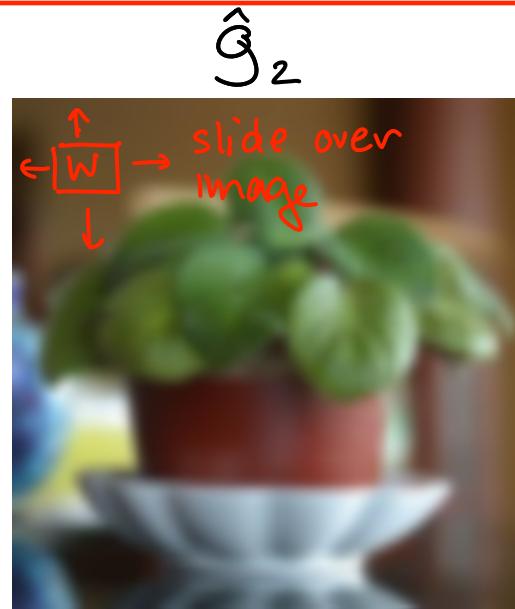


Operation #1: Smooth Image at N-1 Scales

To estimate \hat{g}_3 repeat using \hat{g}_2 as the input:

$$\begin{aligned}\hat{g}_3 &= w * \hat{g}_2 \\ &= \underbrace{(w * w * w)}_{\text{radius is 4 times that of } w} * g_0\end{aligned}$$

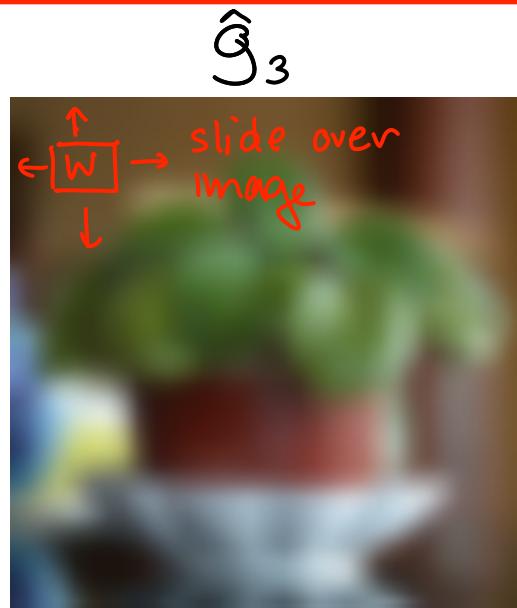
radius is 4 times that of w



Operation #1: Smooth Image at N-1 Scales

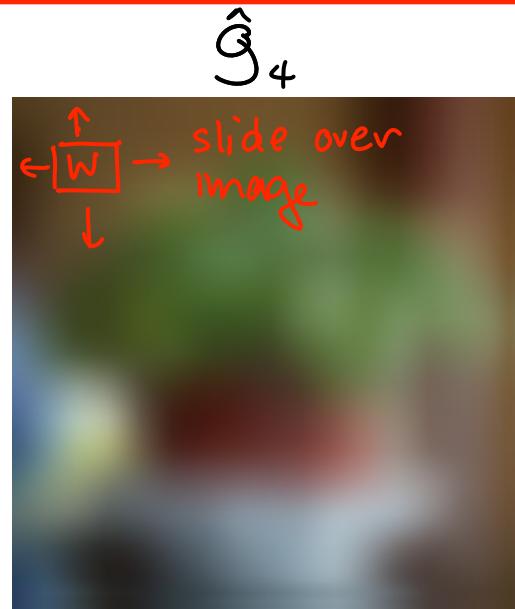
And so on...

$$\begin{aligned}\hat{g}_4 &= w * \hat{g}_3 \\ &= (w * w * w * w) * g_0\end{aligned}$$



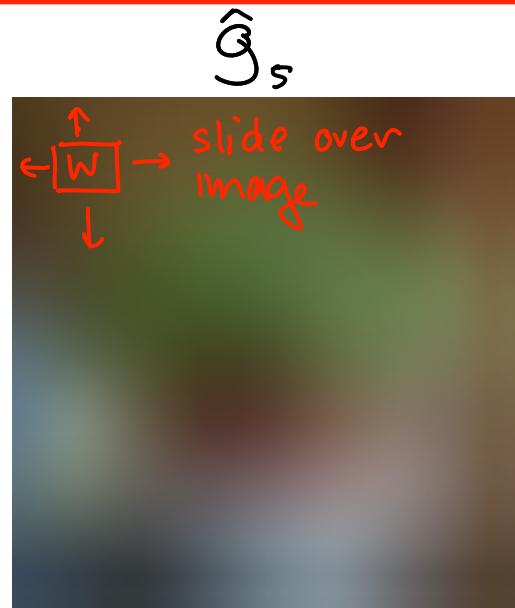
Operation #1: Smooth Image at N-1 Scales

And so on...



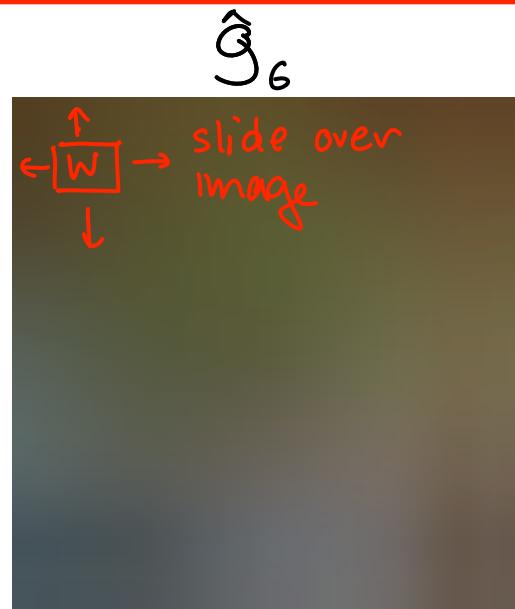
Operation #1: Smooth Image at N-1 Scales

And so on...



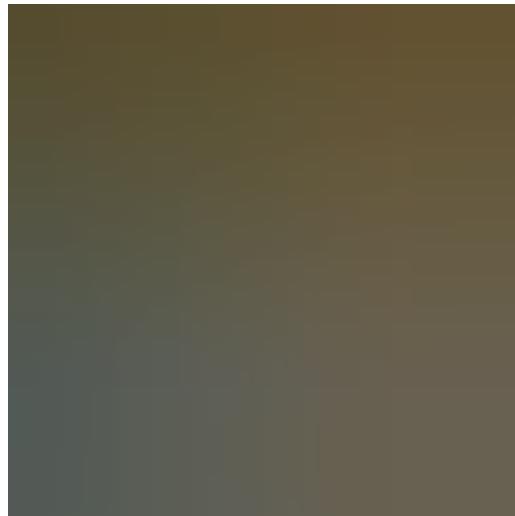
Operation #1: Smooth Image at N-1 Scales

And so on...



Operation #1: Smooth Image at N-1 Scales

$\hat{\mathcal{G}}_7$



2D Kernel to 1D kernel, how?

To take a 5x5 2D kernel into a 1D kernel we must satisfy 4 criteria:

1. The window size must remain the same: 5 elements.
2. The kernel must be symmetric around its origin

$$\hat{w} = \begin{bmatrix} c \\ b \\ a \\ b \\ c \end{bmatrix}$$

3. Applying w to a constant image does not change (intensity preserving)

$$\sum_{m=-2}^2 \hat{w}(m) = 1 \iff a + 2b + 2c = 1$$

4. The ratio of the contributions should follow:

$$a+2c=2b=\frac{1}{2}$$

2D Kernel to 1D kernel, how?

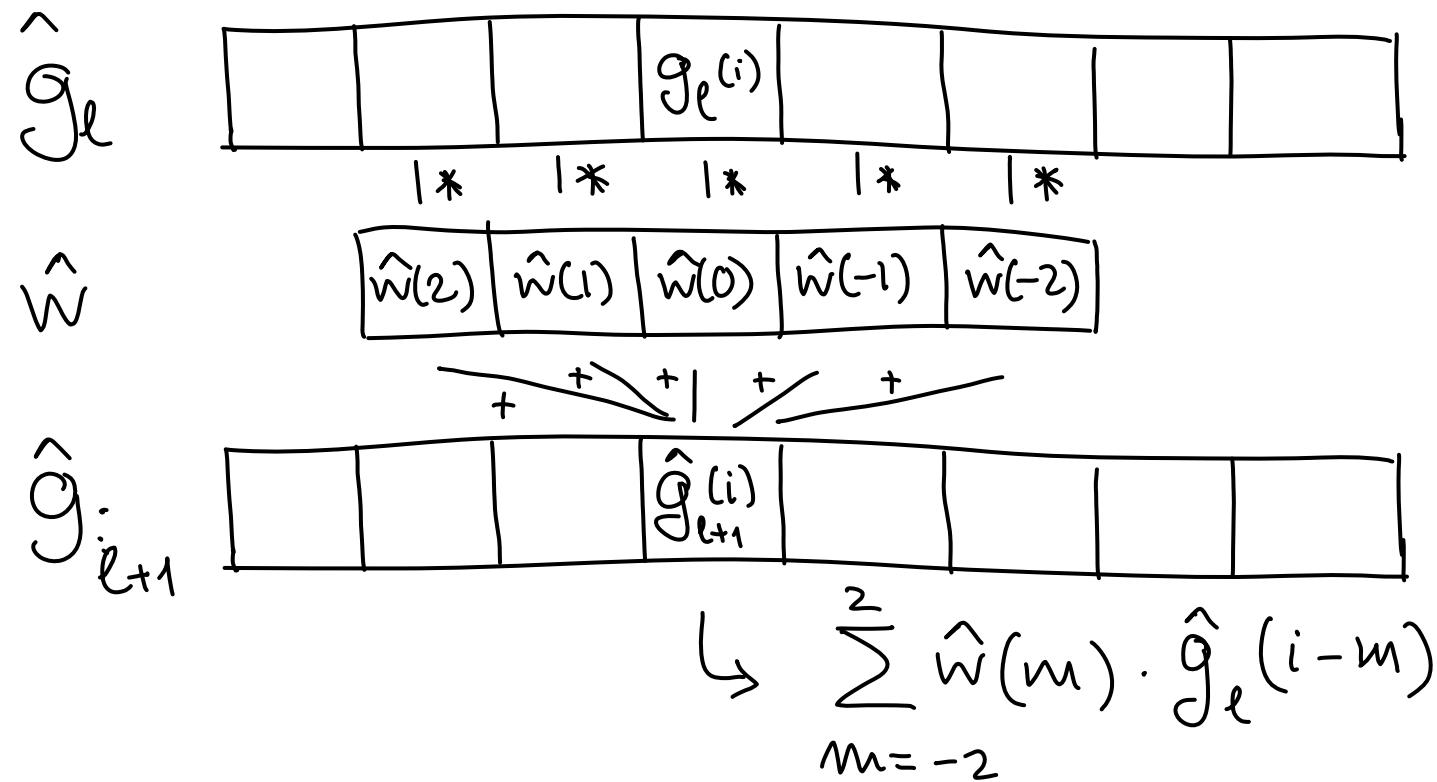
The conditions on the previous slide render filters of the form:

$$\hat{W} = \begin{bmatrix} c \\ b \\ a \\ b \\ c \end{bmatrix} = \begin{bmatrix} \frac{1}{4} - \frac{a}{2} \\ \frac{1}{4} \\ a \\ \frac{1}{4} \\ \frac{1}{4} - \frac{a}{2} \end{bmatrix}$$

usually $a \in [0.3, 0.6]$

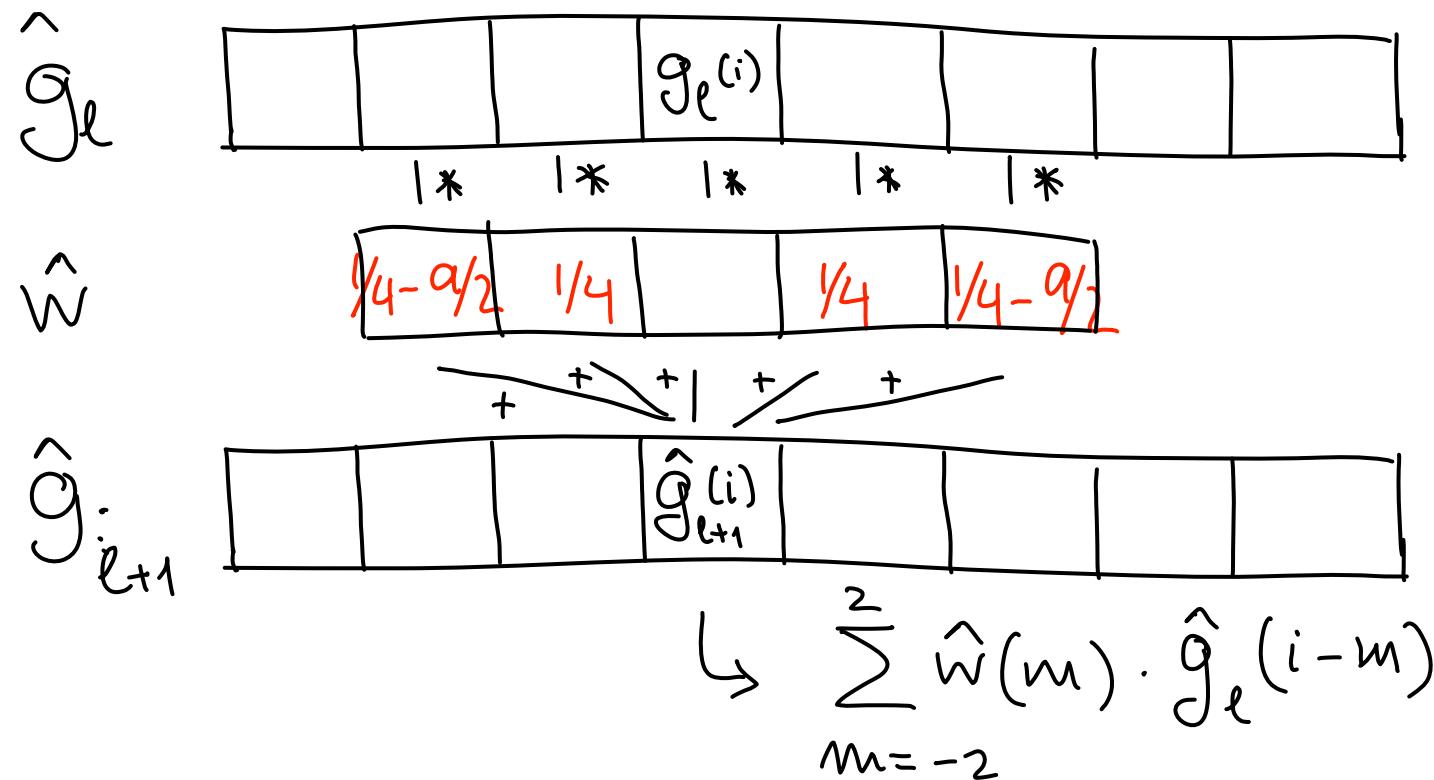
2D Kernel to 1D kernel, how?

This means that to estimate the i^{th} pixel at the $(l+1)^{\text{th}}$ level one must do:



2D Kernel to 1D kernel, how?

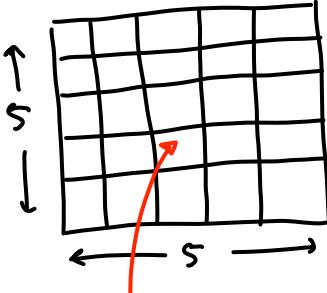
This means that to estimate the i^{th} pixel at the $(l+1)^{\text{th}}$ level one must do:



Defining the Smoothing Filter in 2D

$$\hat{g}_1 = w * g_0$$

w is a separable smoothing filter defined by \hat{w}


$$= \begin{bmatrix} \hat{w} \end{bmatrix} \begin{bmatrix} (\hat{w})^\top \end{bmatrix}$$

\nwarrow 5x1 vector

$$w(m,n) = \hat{w}(m) \cdot \hat{w}(n)$$



Exploiting separability to compute \hat{g}_1 :

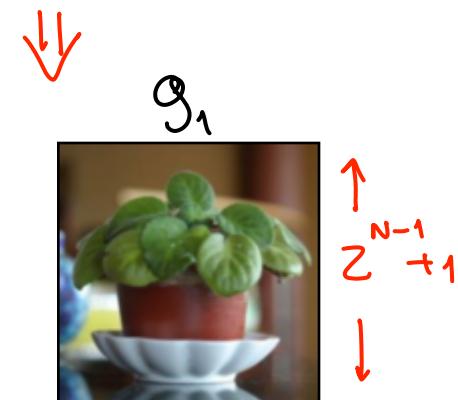
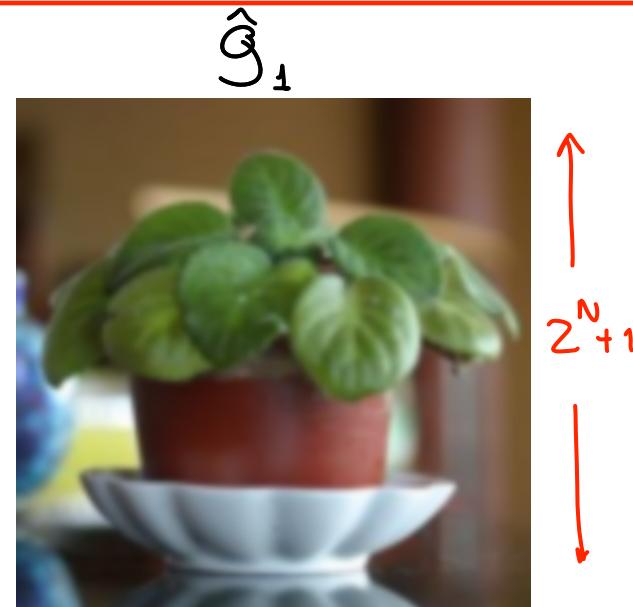
- ① Convolve each row of g_0 with \hat{w}
- ② Convolve the columns of the result with \hat{w} again

Operation #2: Downsample the Smoothed Image

Downsample by taking every other pixel.

Smoothing prevents aliasing effects.

$$g_1(i,j) = \hat{g}_1(2i, 2j)$$



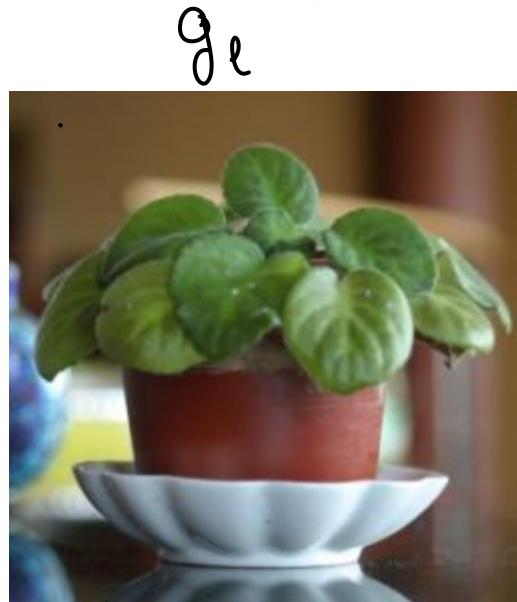
Topic 6.1:

Gaussian & Laplacian Pyramids

- The gaussian pyramid (intro)
- The convolution operation
- Constructing the gaussian pyramid
 - **The REDUCE() function**
- Constructing the Laplacian pyramid
 - The EXPAND() function

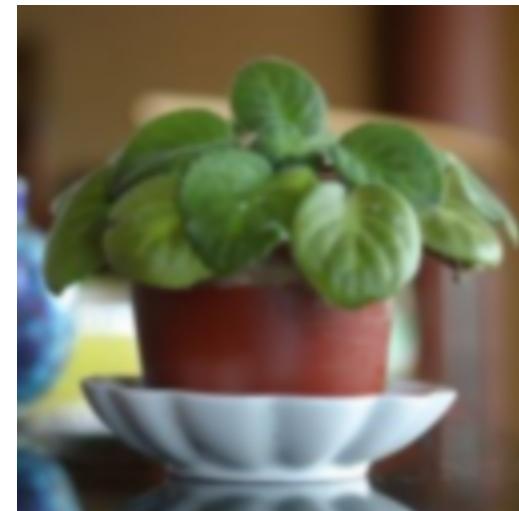
Operations #1 & #2: The REDUCE() Function

The REDUCE() function combines smoothing and down sampling.



$$w * g_e$$

→



$$\leftarrow 2^{N-l} + \rightarrow$$

$$g_{l+1} = \text{REDUCE}(g_e)$$

→

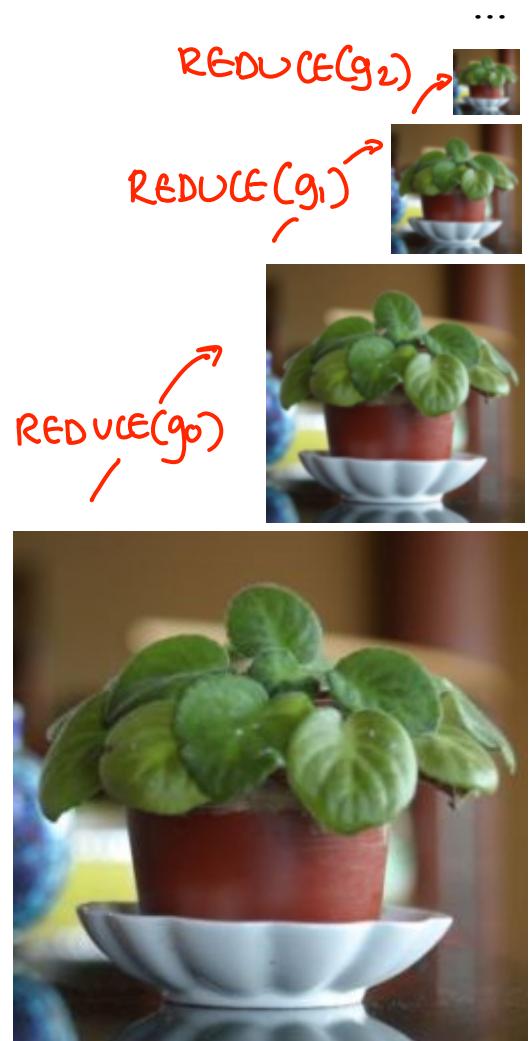
$$g_{l+1}(i,j) = \sum_{m=-2}^2 \sum_{n=-2}^2 w(m,n) \cdot g_e(2i-m, 2j-n)$$



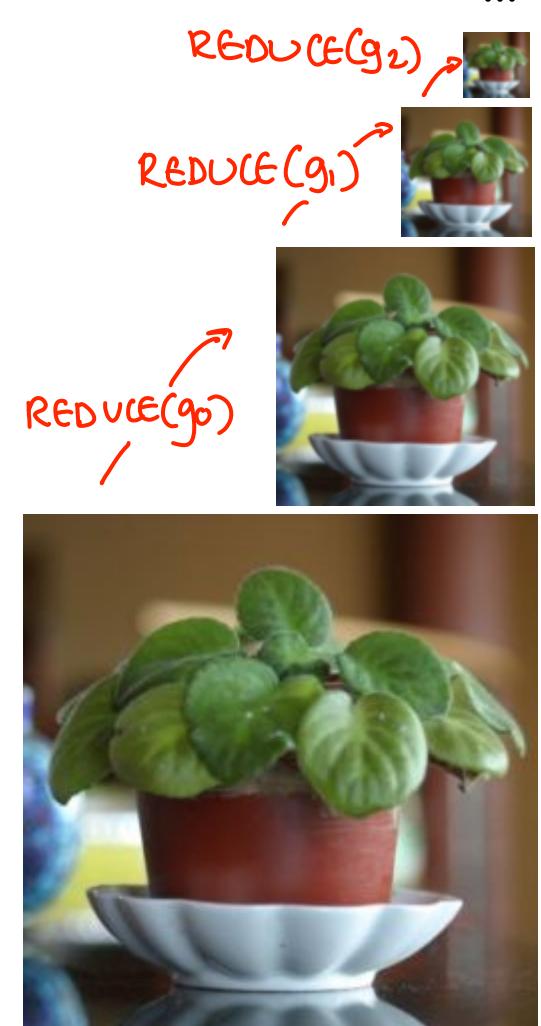
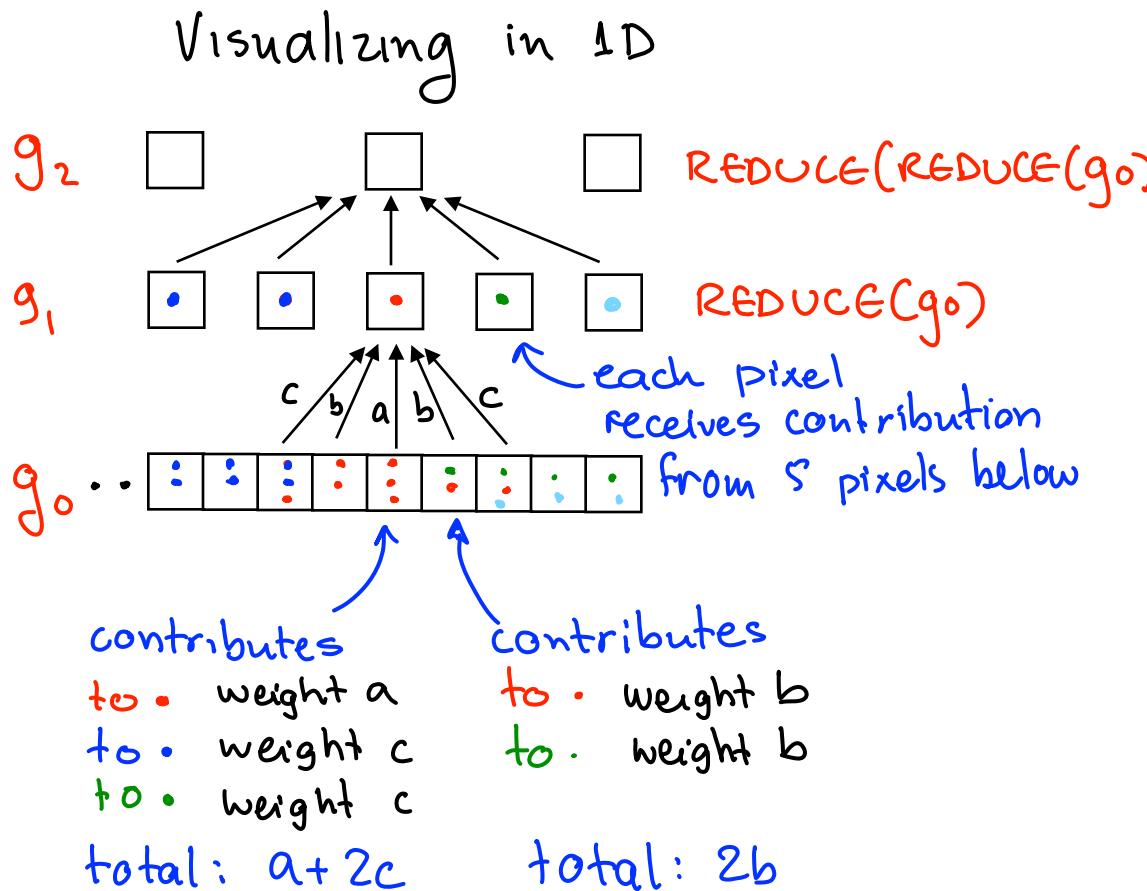
$$\leftarrow 2^{N-l-1} + \rightarrow$$

$$\downarrow \text{down sample } \times 2$$

The REDUCE() function

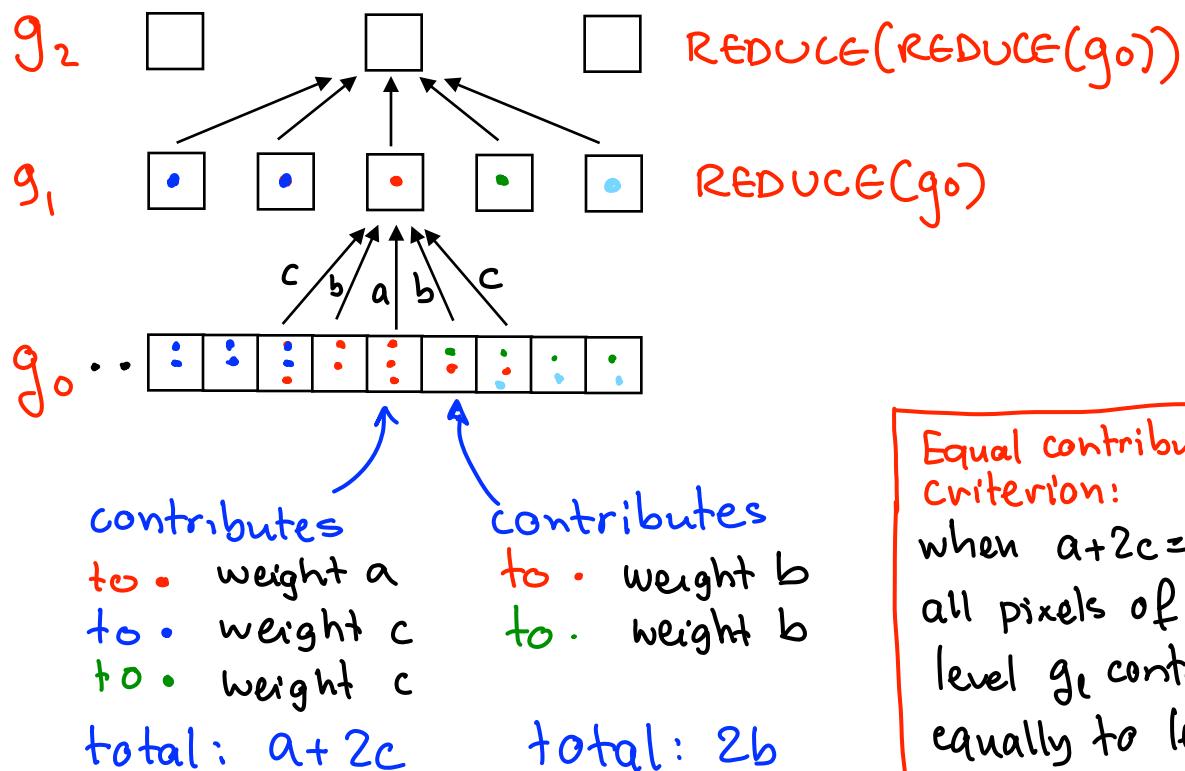


The REDUCE() function

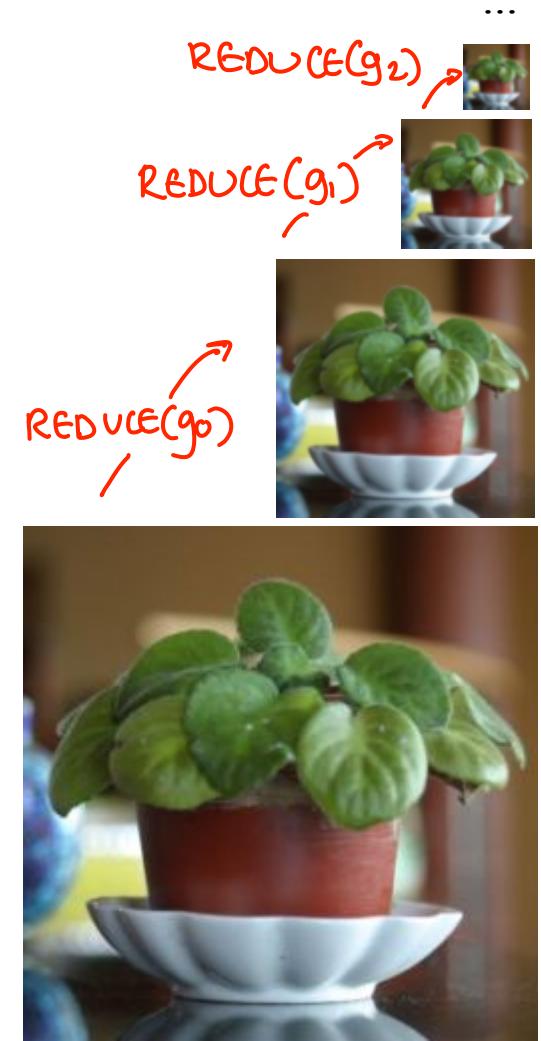


The REDUCE() function

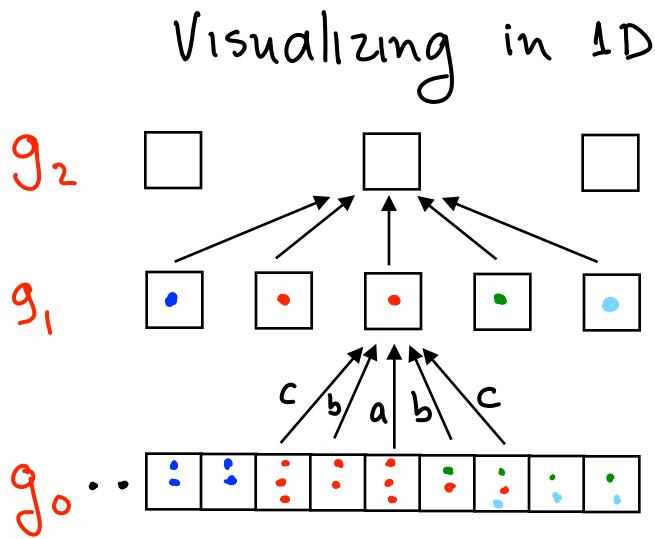
Visualizing in 1D



Equal contribution criterion:
when $a+2c=2b$
all pixels of
level g_1 contribute
equally to level
 g_{l+1}



The REDUCE() function



REDUCE function in matrix notation (1D)

$$g_1 = D_o \cdot C_o \cdot g_0$$

downsampling matrix D_o

convolution matrix C_o

$$g_1 = D_o \cdot C_o \cdot g_0$$

General expression:

$$g_{l+1} = D_l \cdot C_l \cdot g_l$$

What Does Smoothing Take Away?

$g_0 = I$

Original
photo



What Does Smoothing Take Away?

smoothed
photo

$$\hat{g}_1 = W * g_0$$



What Does Smoothing Take Away?

$$g_0 - \hat{g}_1$$



Details in g_0 that were not represented in \hat{g}_1

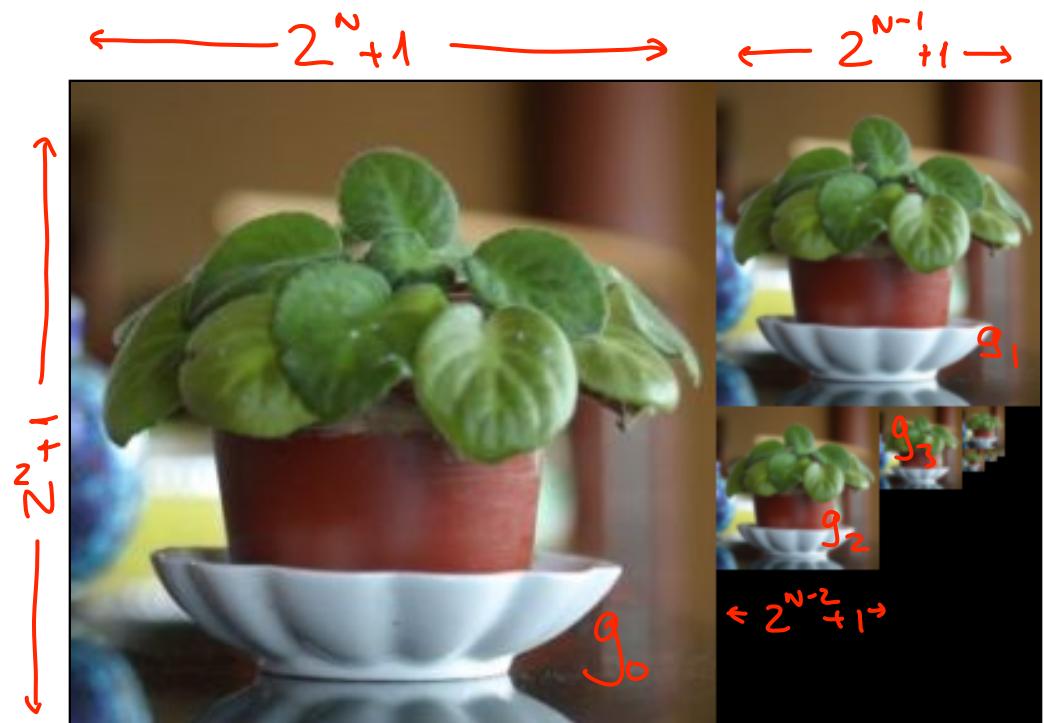
Topic 6.1:

Gaussian & Laplacian Pyramids

- The gaussian pyramid (intro)
- The convolution operation
- Constructing the gaussian pyramid
 - The REDUCE() function
- Constructing the Laplacian pyramid
 - The EXPAND() function

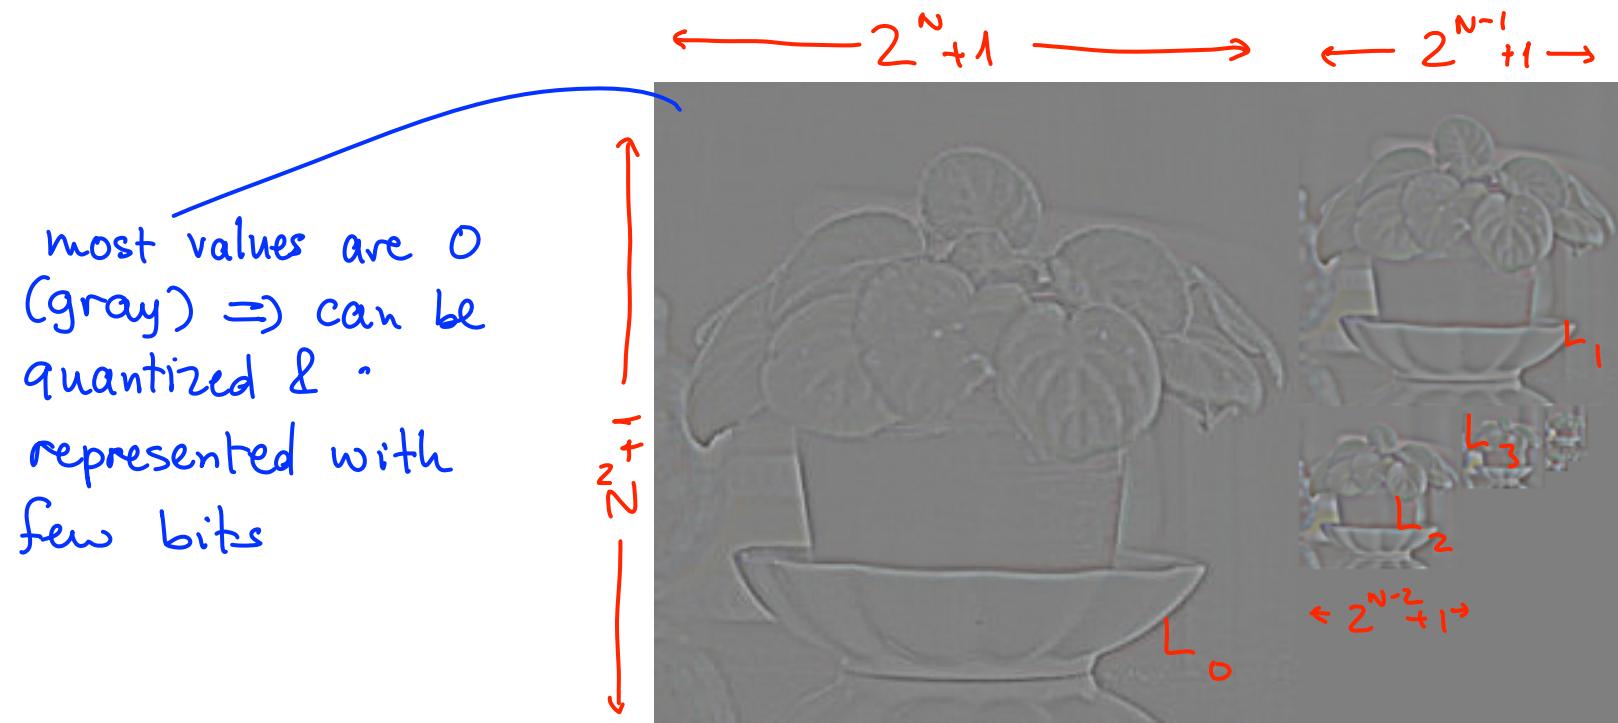
The Laplacian Pyramid

What if instead of storing the smoothed images, we store only the difference between the levels g_l and g_{l+1}



The Laplacian Pyramid

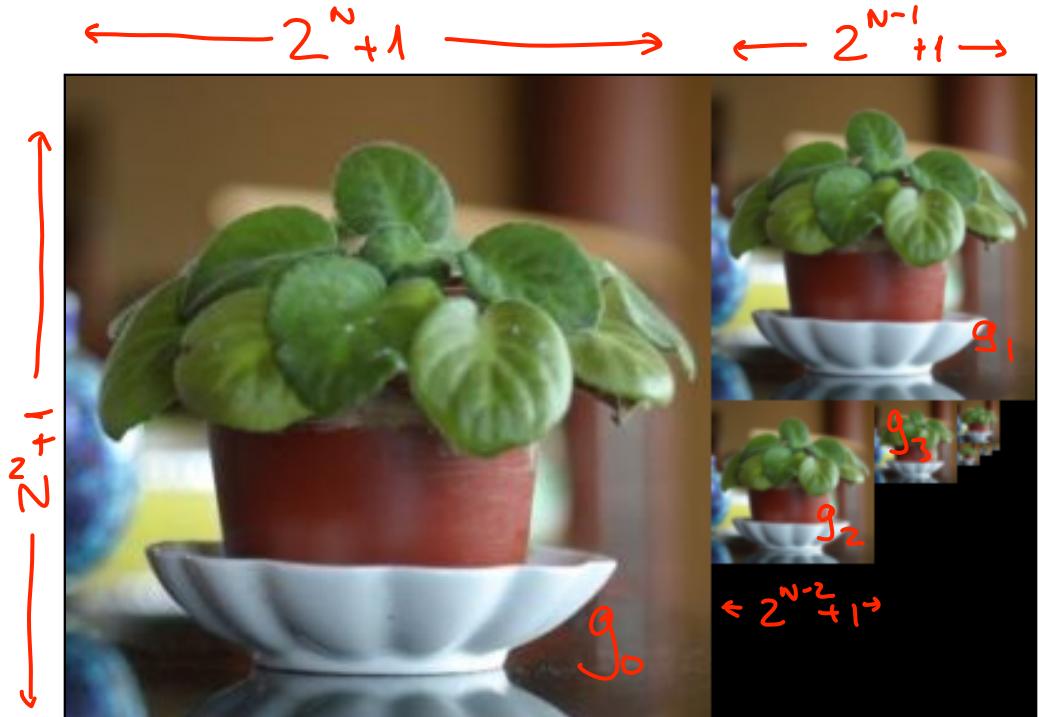
What if instead of storing the smoothed images, we store only the difference between the levels g_i and g_{i+1}



The Laplacian Pyramid

What if instead of storing the smoothed images, we store only the difference between the levels g_l and g_{l+1}

- But g_{l+1}, g_l are not the same size!
⇒ Expand g_{l+1} to make it equal size to g_l



Operation #3: The EXPAND() Function

And we can write this with an equation:

$$\leftarrow 2^{N-l} + 1 \rightarrow$$



EXPAND()



g_l



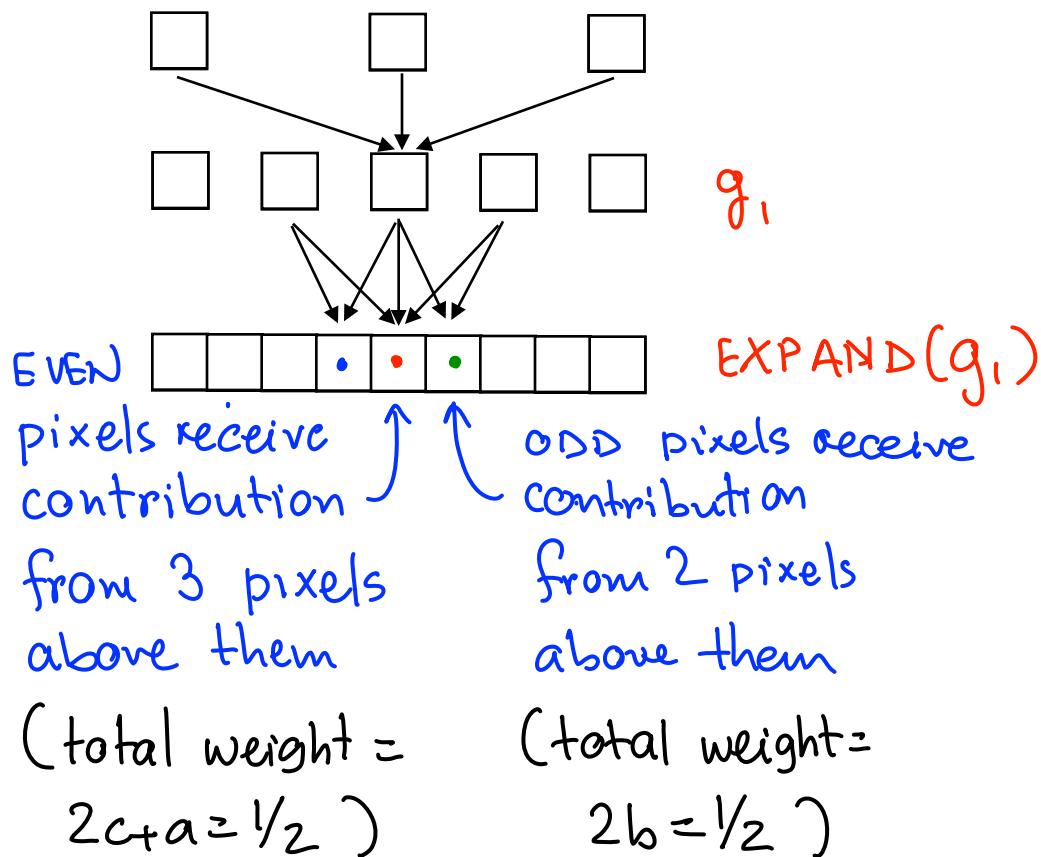
EXPAND(g_l)

$$\boxed{\text{EXPAND}(g_l) = 4 \sum_{m=-2}^2 \sum_{n=-2}^2 w(m,n) g_l\left(\frac{i-m}{2}, \frac{j-n}{2}\right)}$$

The EXPAND() function

- The function upsamples level g_e by doubling its size from $(2^{N-l+1}) \times (2^{N-l+1})$ to $(2^{N-l+1}+1) \times (2^{N-l+1}+1)$

Visualizing in 1D



General expression in 1D

$$\text{EXPAND}(g_e)(i) = 2 \sum_{m=-2}^2 \hat{w}(m) \cdot g_e\left(\frac{i-m}{2}\right)$$

evaluated only for i, m where $\frac{i-m}{2}$ is an integer

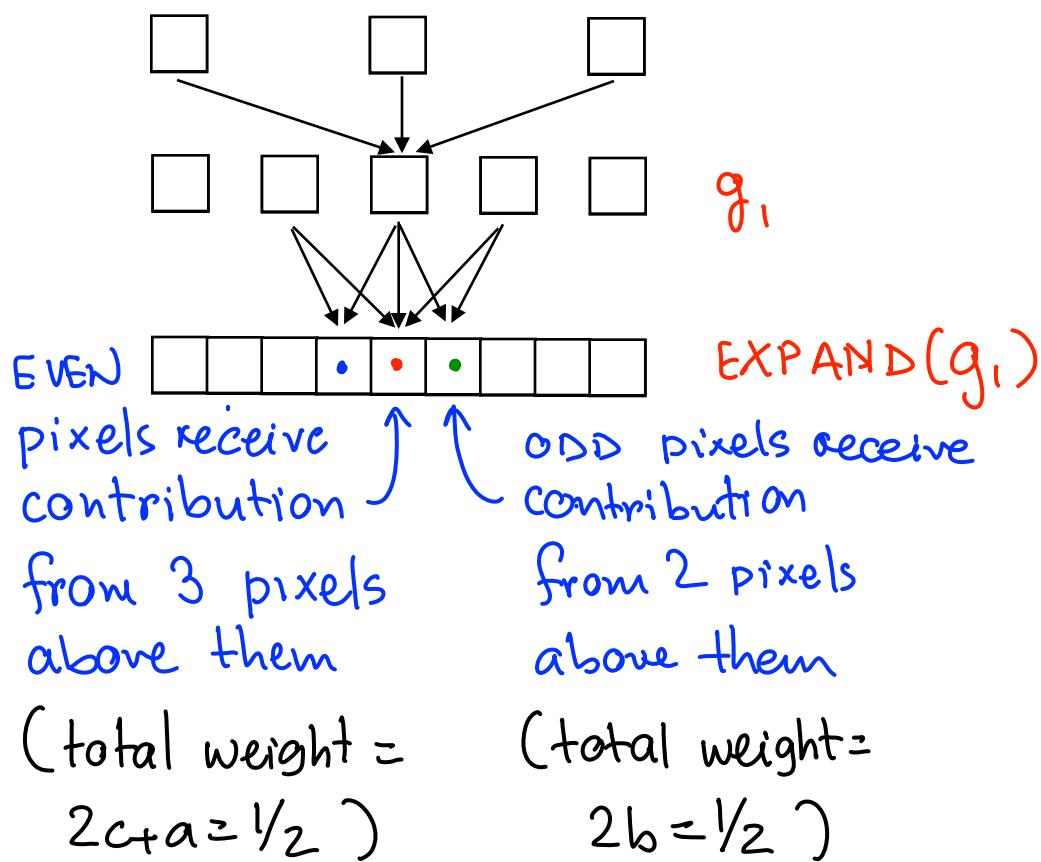
The EXPAND() function

- The function upsample^s level g_e by doubling its size from $(2^{N-l+1}) \times (2^{N-l}+1)$ to $(2^{N-l+1}) \times (2^{N-l+1}+1)$

The EXPAND() function

- The function upsamples level g_e by doubling its size from $(2^{N-l}+1) \times (2^{N-l}+1)$ to $(2^{N-l+1}+1) \times (2^{N-l+1}+1)$

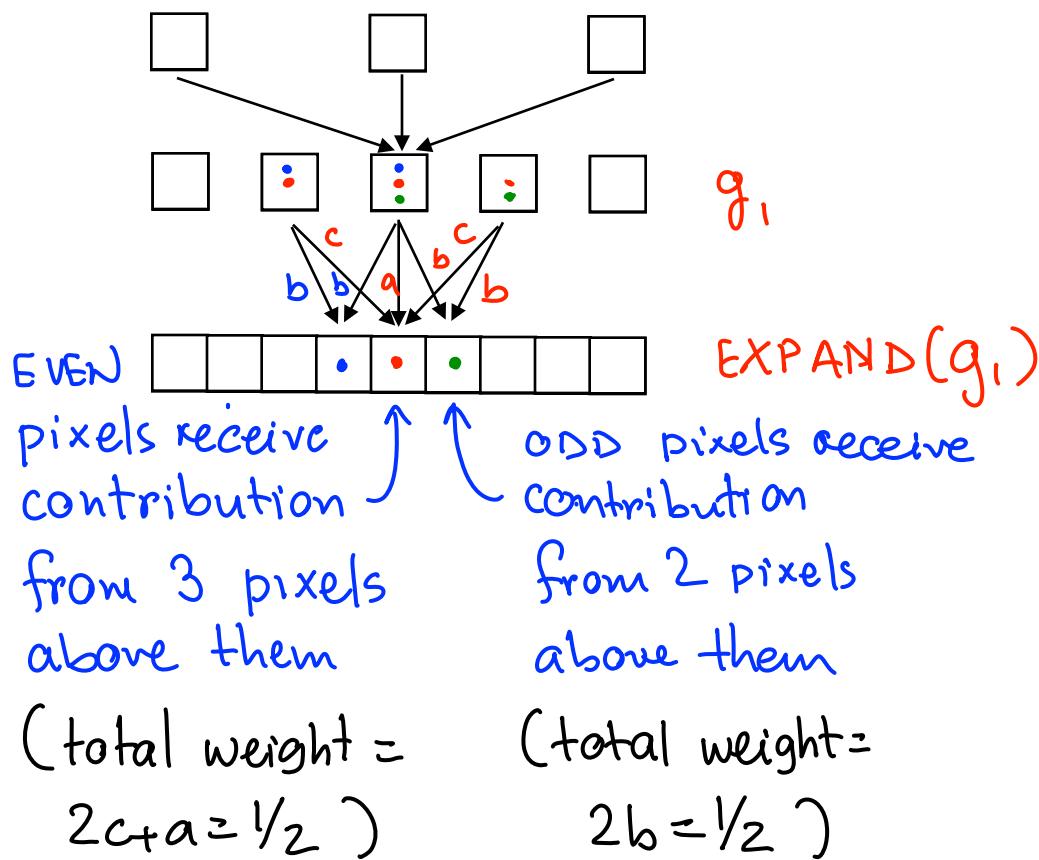
Visualizing in 1D



The EXPAND() function

- The function upsamples level g_e by doubling its size from $(2^{N-l+1}) \times (2^{N-l+1})$ to $(2^{N-l+1}) \times (2^{N-l+1}+1)$

Visualizing in 1D



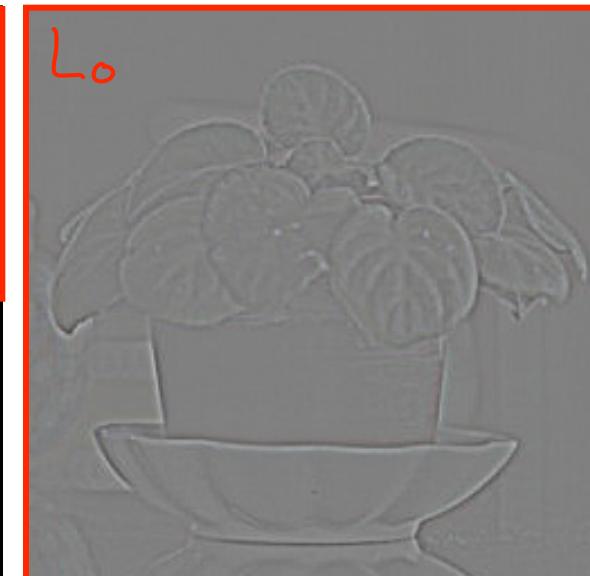
EXPAND function in matrix notation (1D)

$$\text{EXPAND}(g_1) = \begin{bmatrix} 2^{N+1} & \uparrow \\ \text{EXPAND}(g_1) & \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & \cdots & 0 \\ c & b & a & b & c \\ c & b & a & b & c \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & \vdots & \vdots & \ddots & \vdots \end{bmatrix} \begin{bmatrix} g_1 \\ \text{upsampling convolution matrix } C_0 \\ \text{matrix } D_0^T \end{bmatrix}$$

$$\boxed{\text{EXPAND}(g_1) = C_0 \cdot (D_0)^T g_1}$$

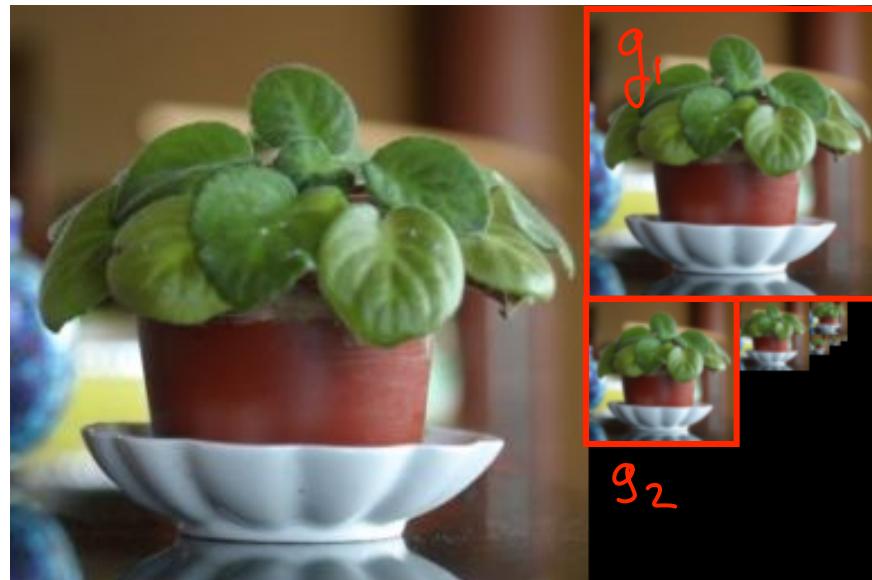
The Laplacian Pyramid

$$L_0 = g_0 - \text{EXPAND}(g_1)$$



The Laplacian Pyramid

$$L_1 = g_1 - \text{EXPAND}(g_2)$$



The Laplacian Pyramid

$$L_2 = g_2 - \text{EXPAND}(g_3)$$



The Laplacian Pyramid

- Often we use a truncated Laplacian pyramid by storing images $L_0, L_1, \dots, L_k, g_{k+1}$ for $k+1 < N$

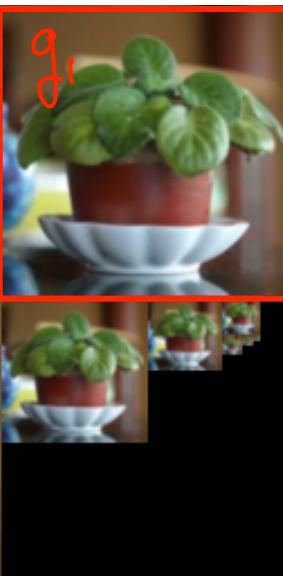
Base case: store g_N .



The Laplacian Pyramid

But we can use the Pyramid in the opposite way too. To recover g_0 , we can do:

$$g_0 = L_0 + \text{EXPAND}(g_1)$$

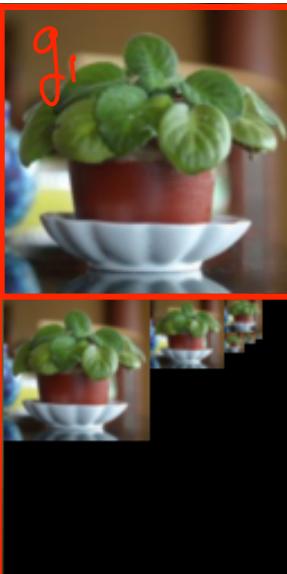
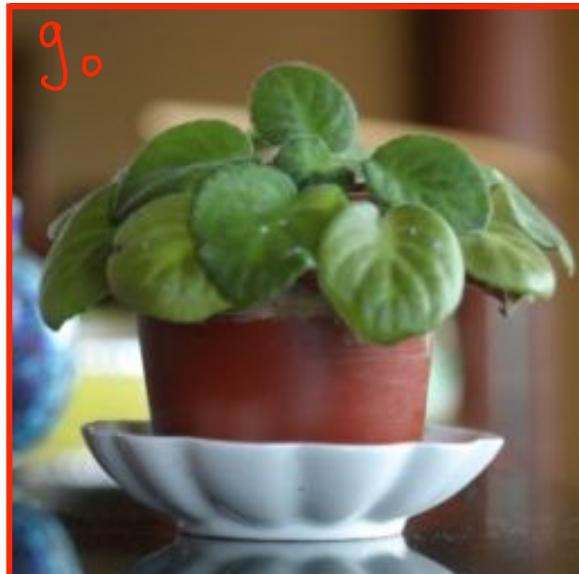


The Laplacian Pyramid

But we can use the Pyramid in the opposite way too. To recover g_0 , we can do: $g_0 = L_0 + \text{EXPAND}(g_1)$

Given an input image of size $(2^N+1) \times (2^N+1)$,
the Laplacian pyramid representation consists of

- L_0, \dots, L_{K-1}
- g_k for some $K \leq N$



Transmission using EXPAND

RECEIVER

TRANSMITTER

③ Display :

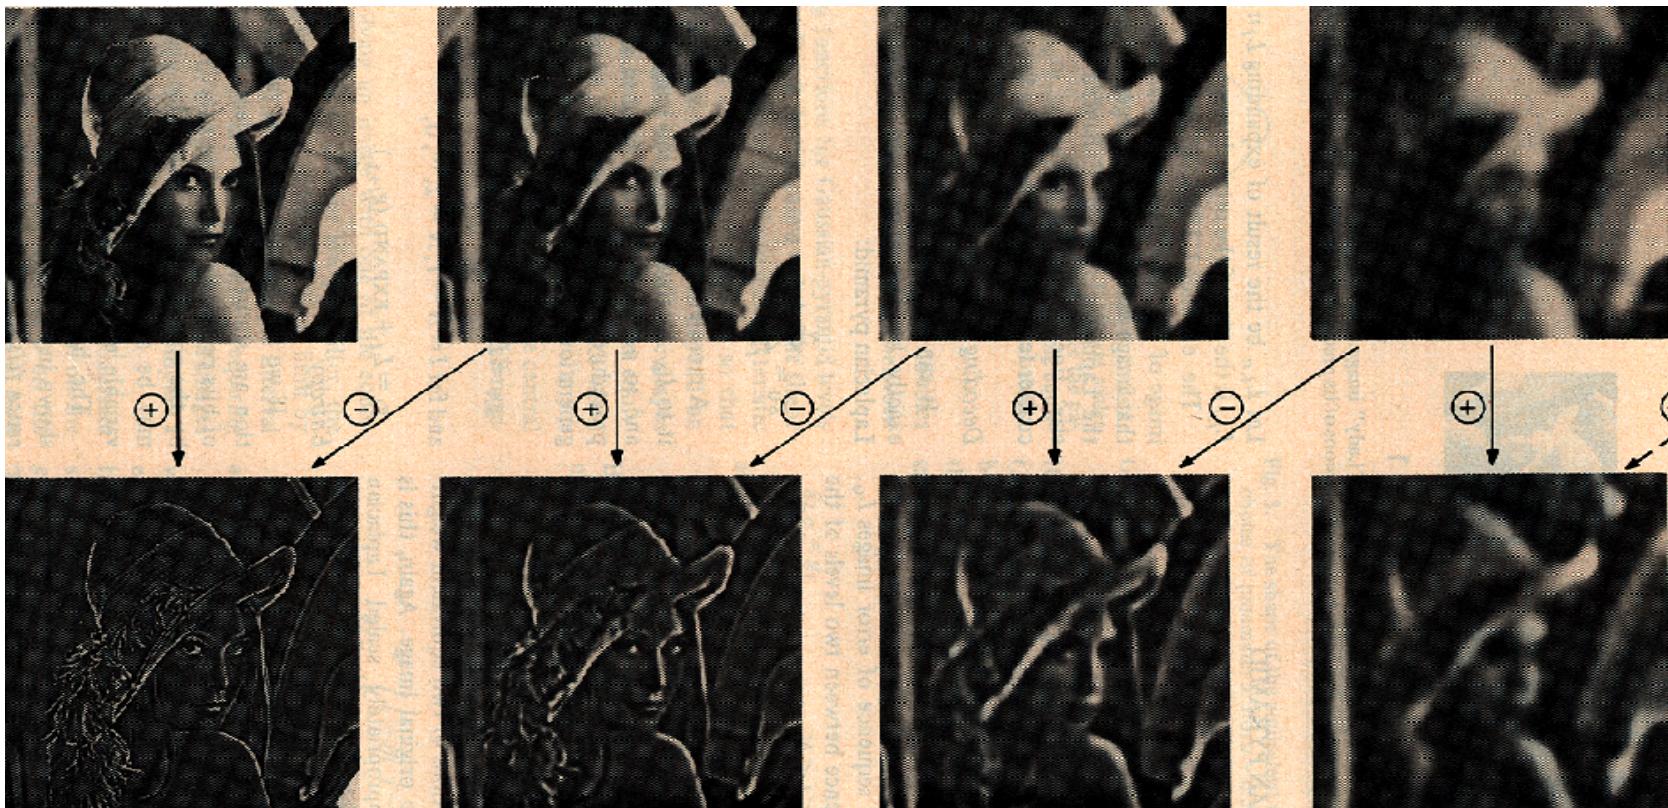
$$d_3 = d_2 + \\ \text{EXPAND}^{k-2}(L_{k-2})$$

② Display

$$d_2 = d_1 + \\ \text{EXPAND}^{k-1}(L_{k-1})$$

① Display $d_1 = \text{EXPAND}^k(g_k)$

① Transmit g_k



⑤ Transmit
 L_{k-4}

④ Transmit
 L_{k-3}

③ Transmit
 L_{k-2}

② Transmit
 L_{k-1}

Topic 6:

Hierarchical image representations

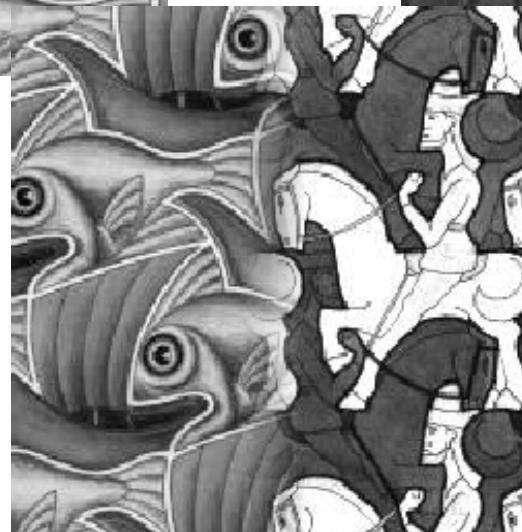
1. Gaussian & Laplacian pyramids
2. Applications:
 1. Multi-resolution image blending
 2. Multi-resolution image editing
 3. Multi-resolution texture synthesis

Image Blending

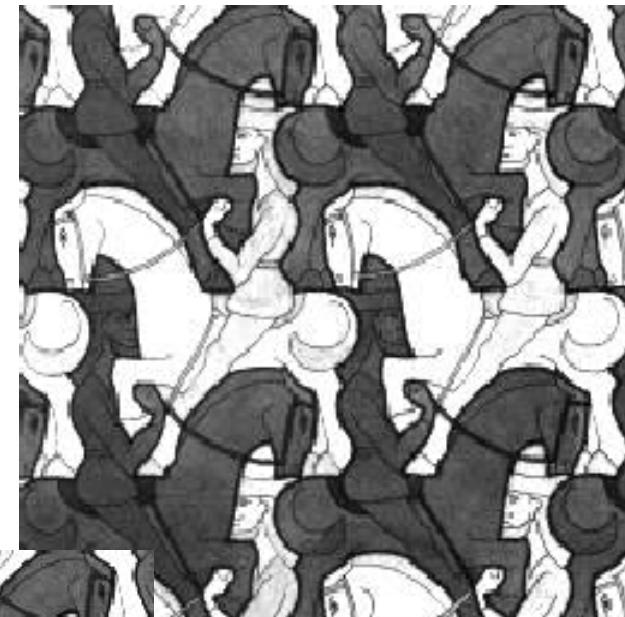
Source 1



Blend



Source 2

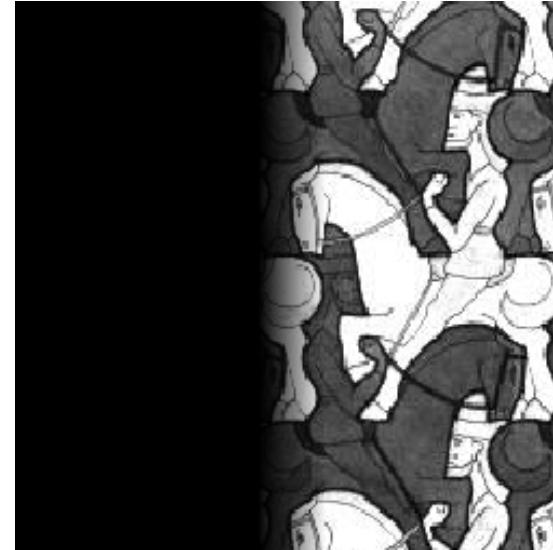


Slides adapted
from A. Efros
(CMU)

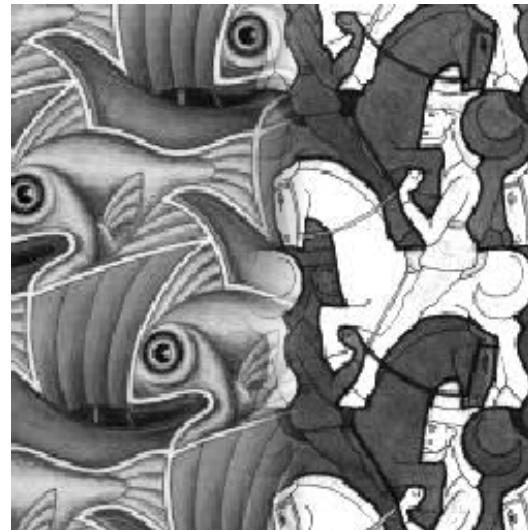
Feathering



+

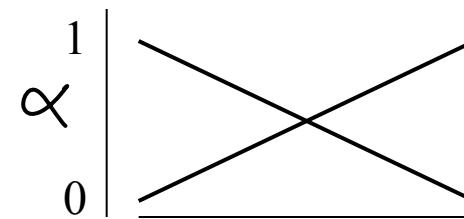
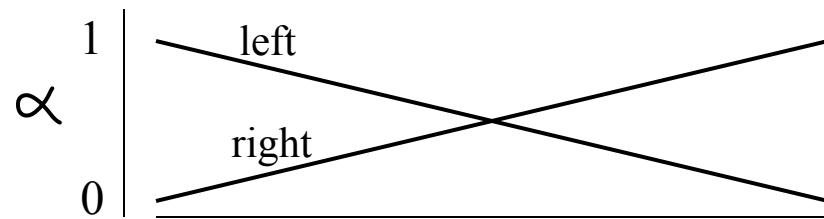
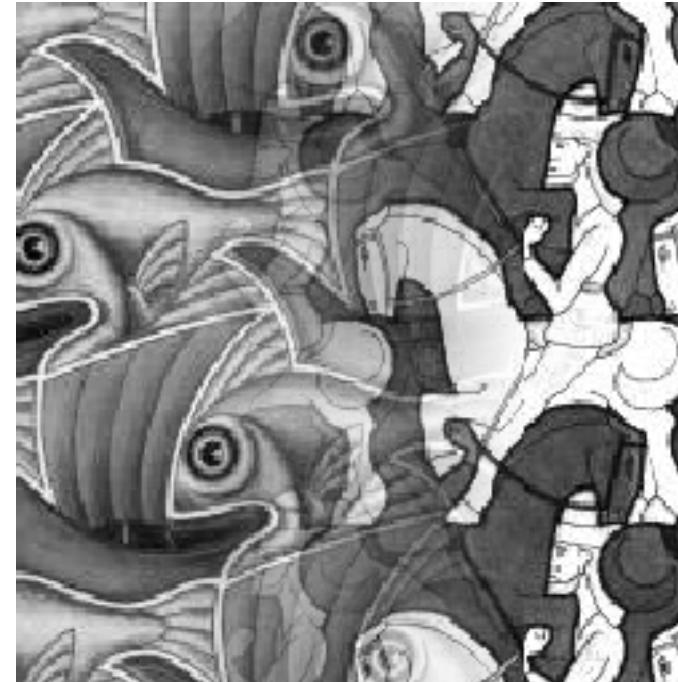
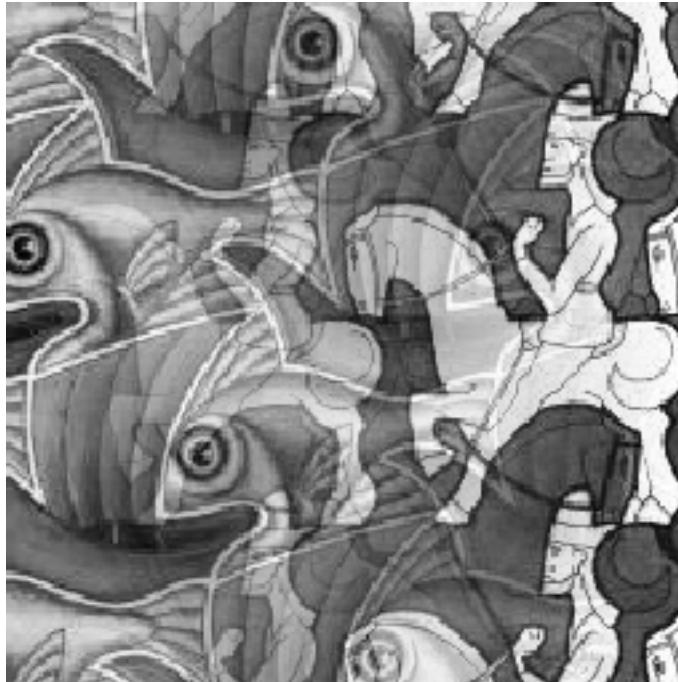


=



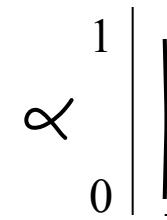
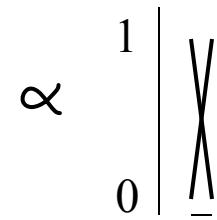
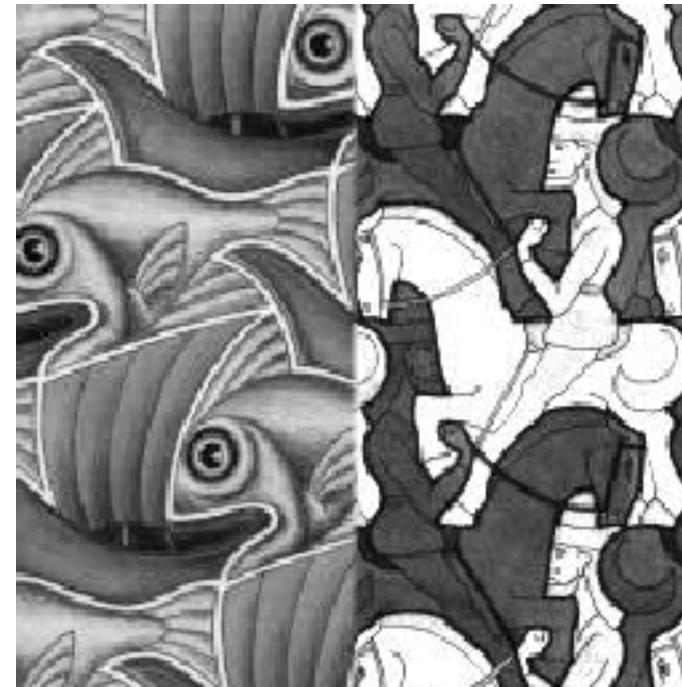
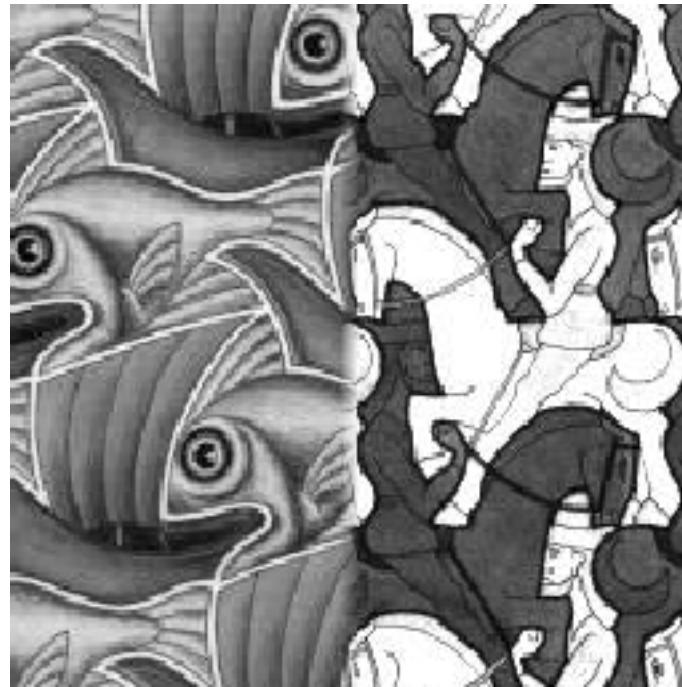
Assign an alpha value
to each pixel near
the seam to make it
less visible

Effect of Window Size



Ghosting is most
apparent here

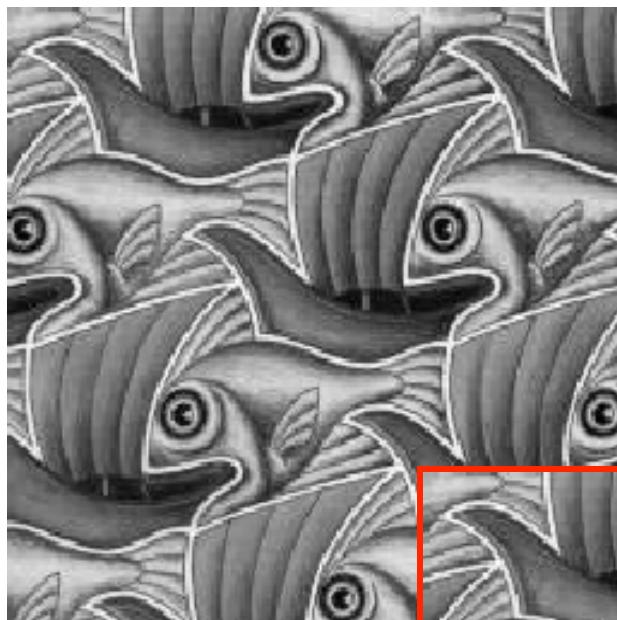
Effect of Window Size



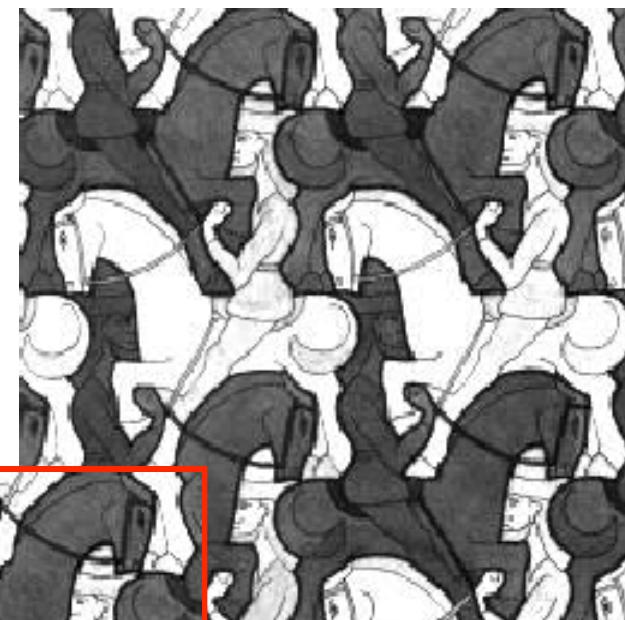
seam most visible
here

Image Blending

Source A



Source B



Blend



Main challenge:

- Minimize ghosting without making visible seams

Slides adapted
from A. Efros
(CMU)

Application #1: Pyramid Blending Algorithm

Input: Source images A, B & binary matte M
Output: Blended image S

A



B



Pyramid Blending Algorithm

Input: Source images A, B & binary matte M.
Output: Blended image S

A



B

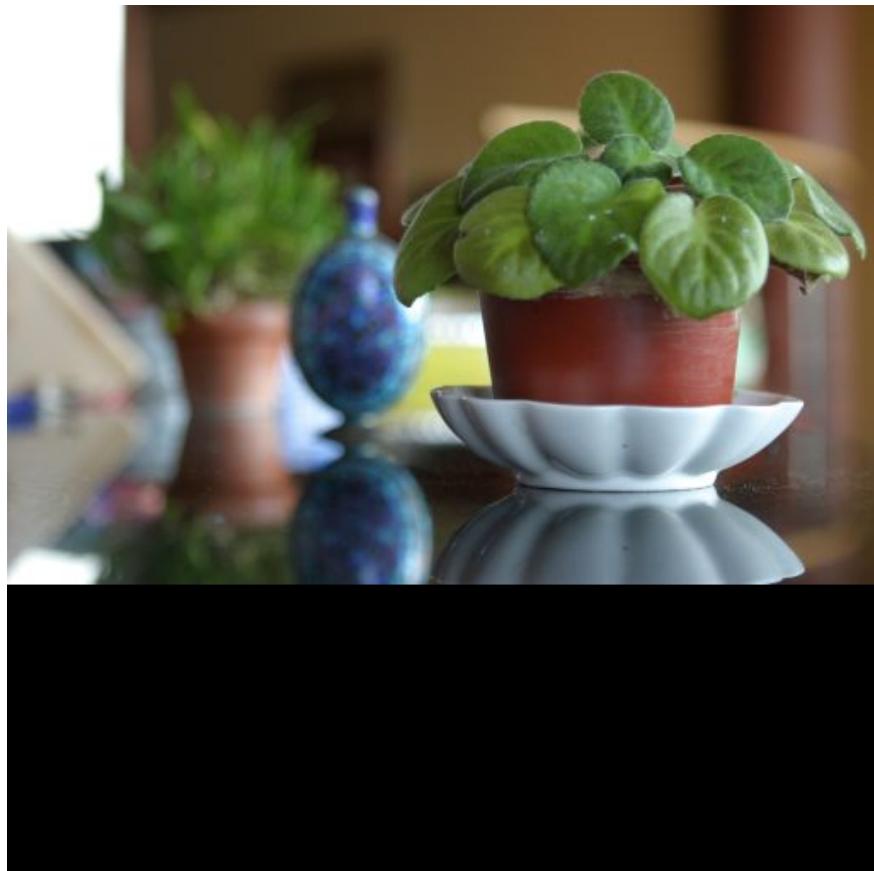


Pad to make
size $(2^N+1) \times (2^N+1)$

Pyramid Blending Algorithm

Input: Source images A, B & binary matte M.
Output: Blended image S

A

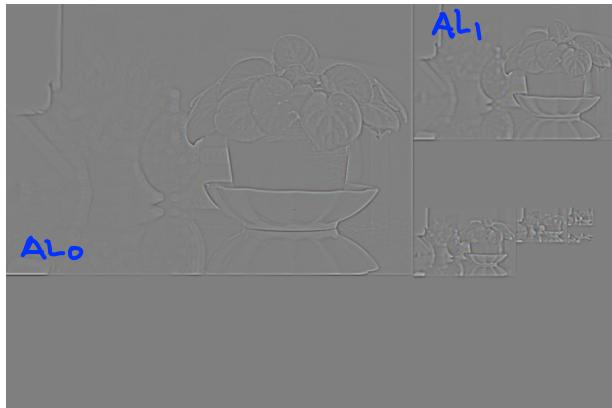


M



Pyramid Blending Algorithm

Input: Source images A, B & binary matte M
Output: Blended image S



① Compute A's

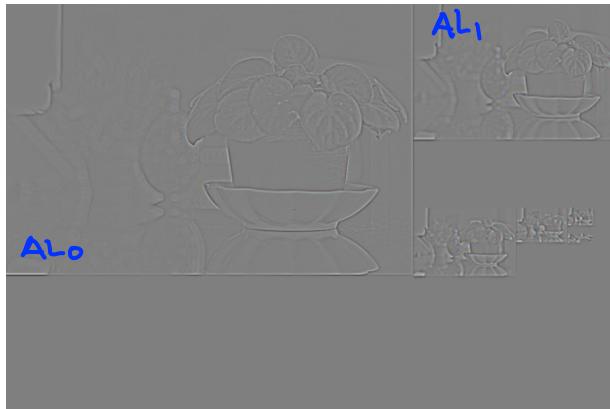
Laplacian

pyramid:

$AL_0, \dots, AL_{N-1}, AG_N$

Pyramid Blending Algorithm

Input: Source images A, B & binary matte M
Output: Blended image S

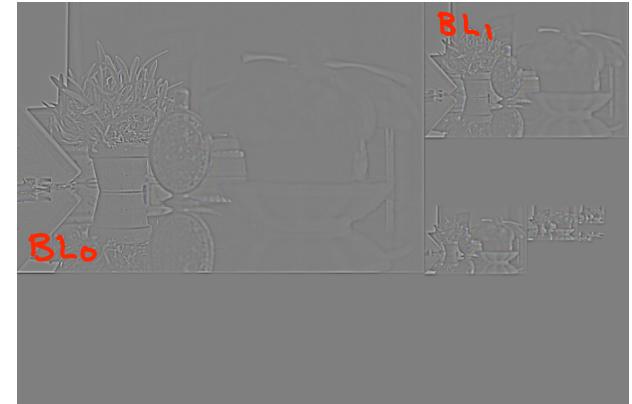
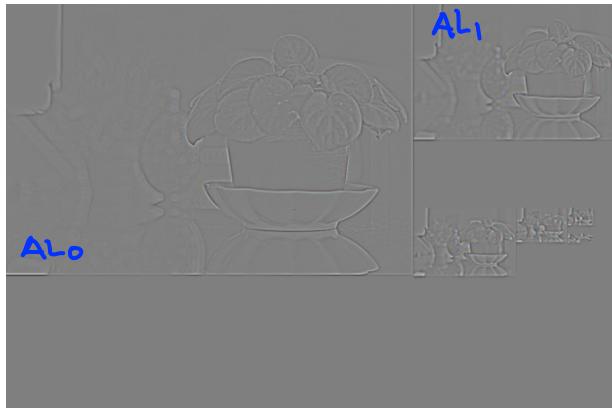


① Compute A's
Laplacian
pyramid:
 $AL_0, \dots, AL_{N-1}, AG_N$

② Compute B's
Laplacian
pyramid:
 $BL_0, \dots, BL_{N-1}, BG_N$

Pyramid Blending Algorithm

Input: Source images A, B & binary matte M
Output: Blended image S



① Compute A's
Laplacian
pyramid:
 $AL_0, \dots, AL_{N-1}, AG_N$

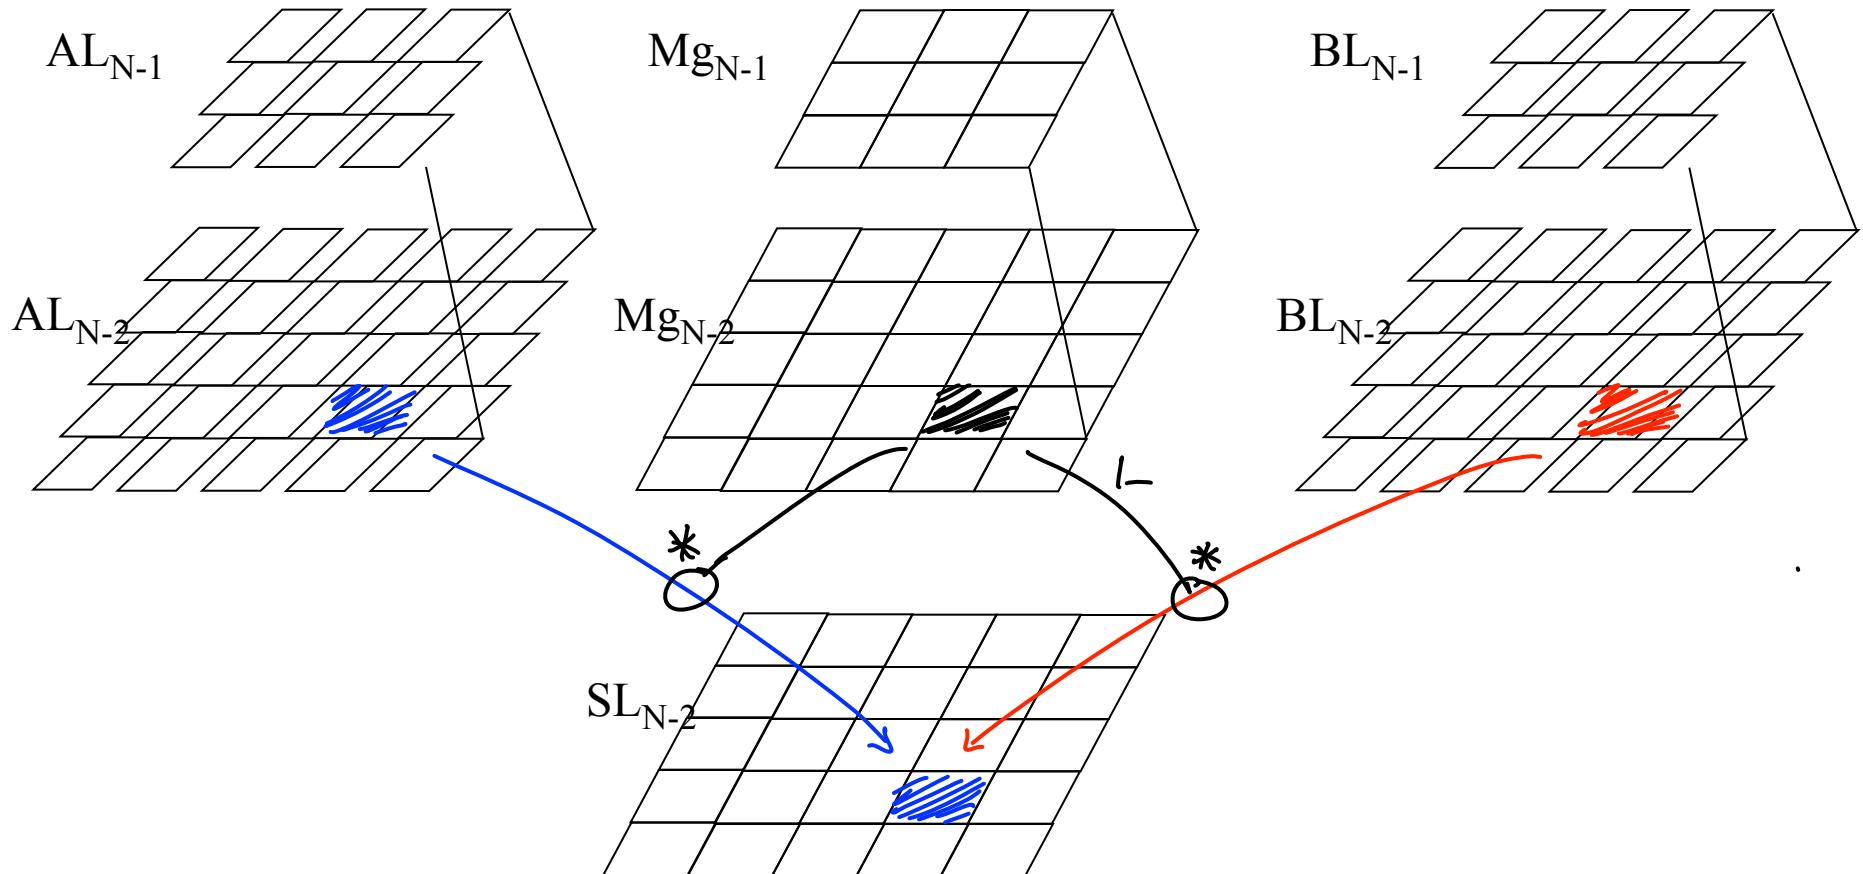
③ Compute M's
Gaussian
pyramid:
 Mg_0, \dots, Mg_N

② Compute B's
Laplacian
pyramid:
 $BL_0, \dots, BL_{N-1}, BG_N$

Pyramid Blending Algorithm

④ Compute the Laplacian pyramid, $SL_0, SL_1, \dots, SL_{N-1}, SG_N$ by applying the matting equation with matte Mg_e :

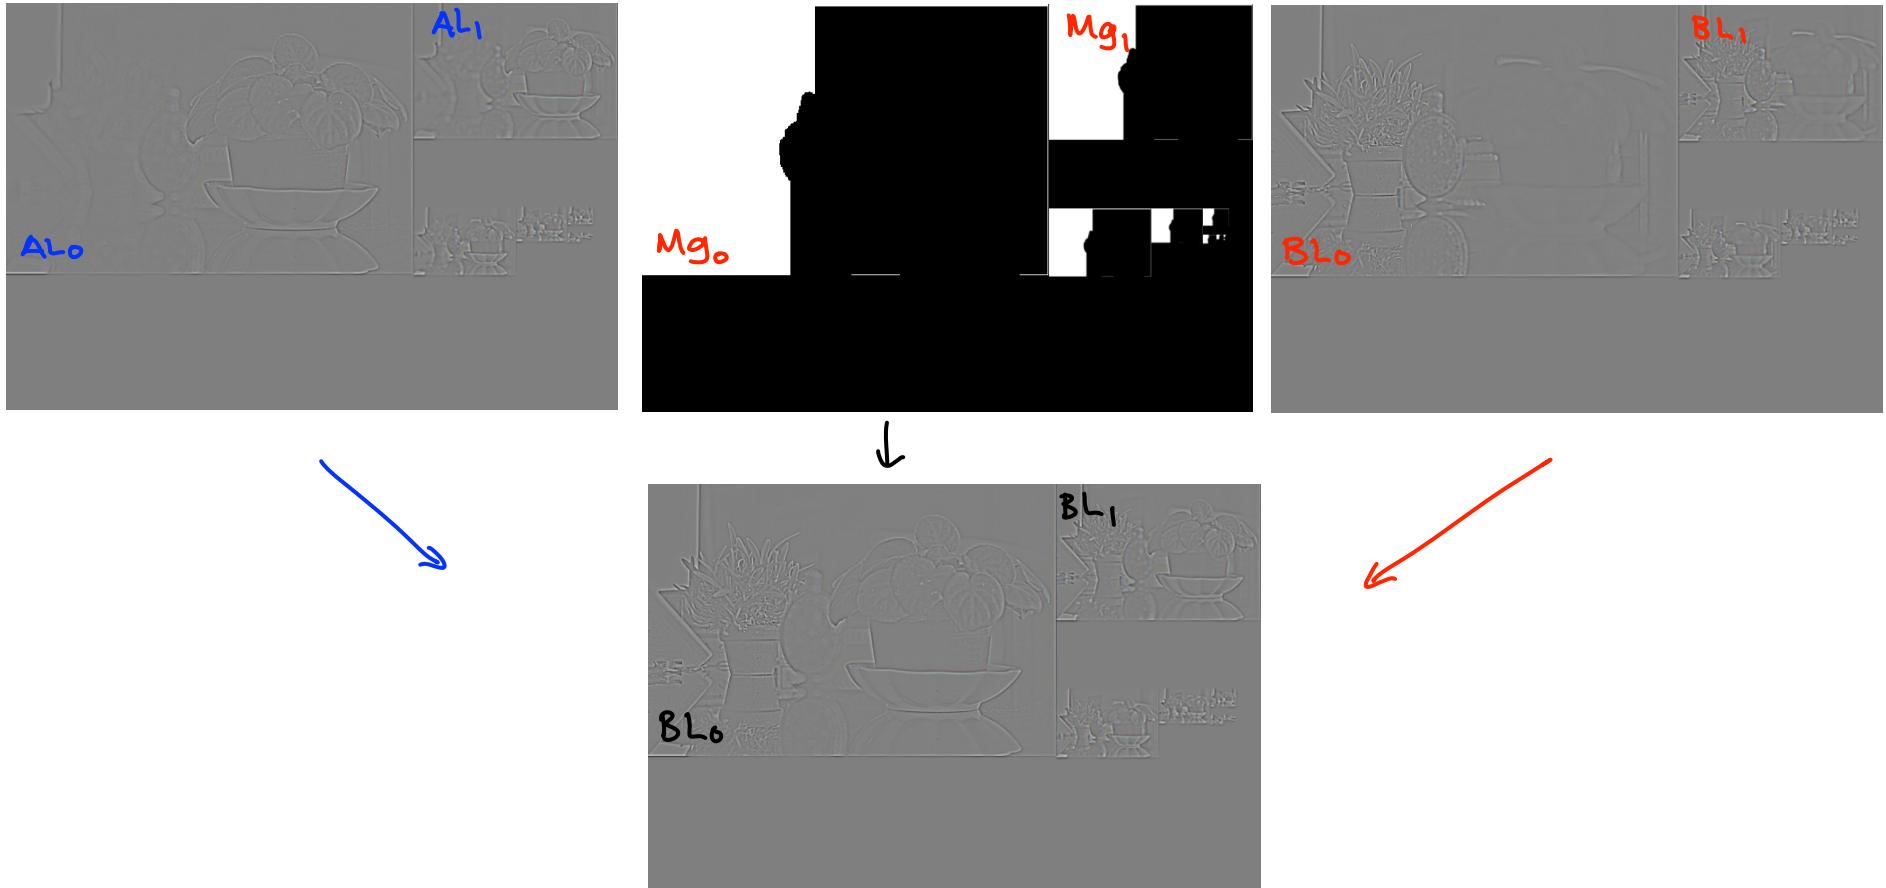
$$SL_\ell(i,j) = Mg_e(i,j) AL_\ell(i,j) + (1 - Mg_e(i,j)) BL_\ell(i,j)$$



Pyramid Blending Algorithm

④ Compute the Laplacian pyramid, $SL_0, SL_1, \dots, SL_{n-1}, SG_n$ by applying the matting equation with matte Mg_e :

$$SL_\ell(i,j) = Mg_e(i,j) AL_\ell(i,j) + (1 - Mg_e(i,j)) BL_\ell(i,j)$$

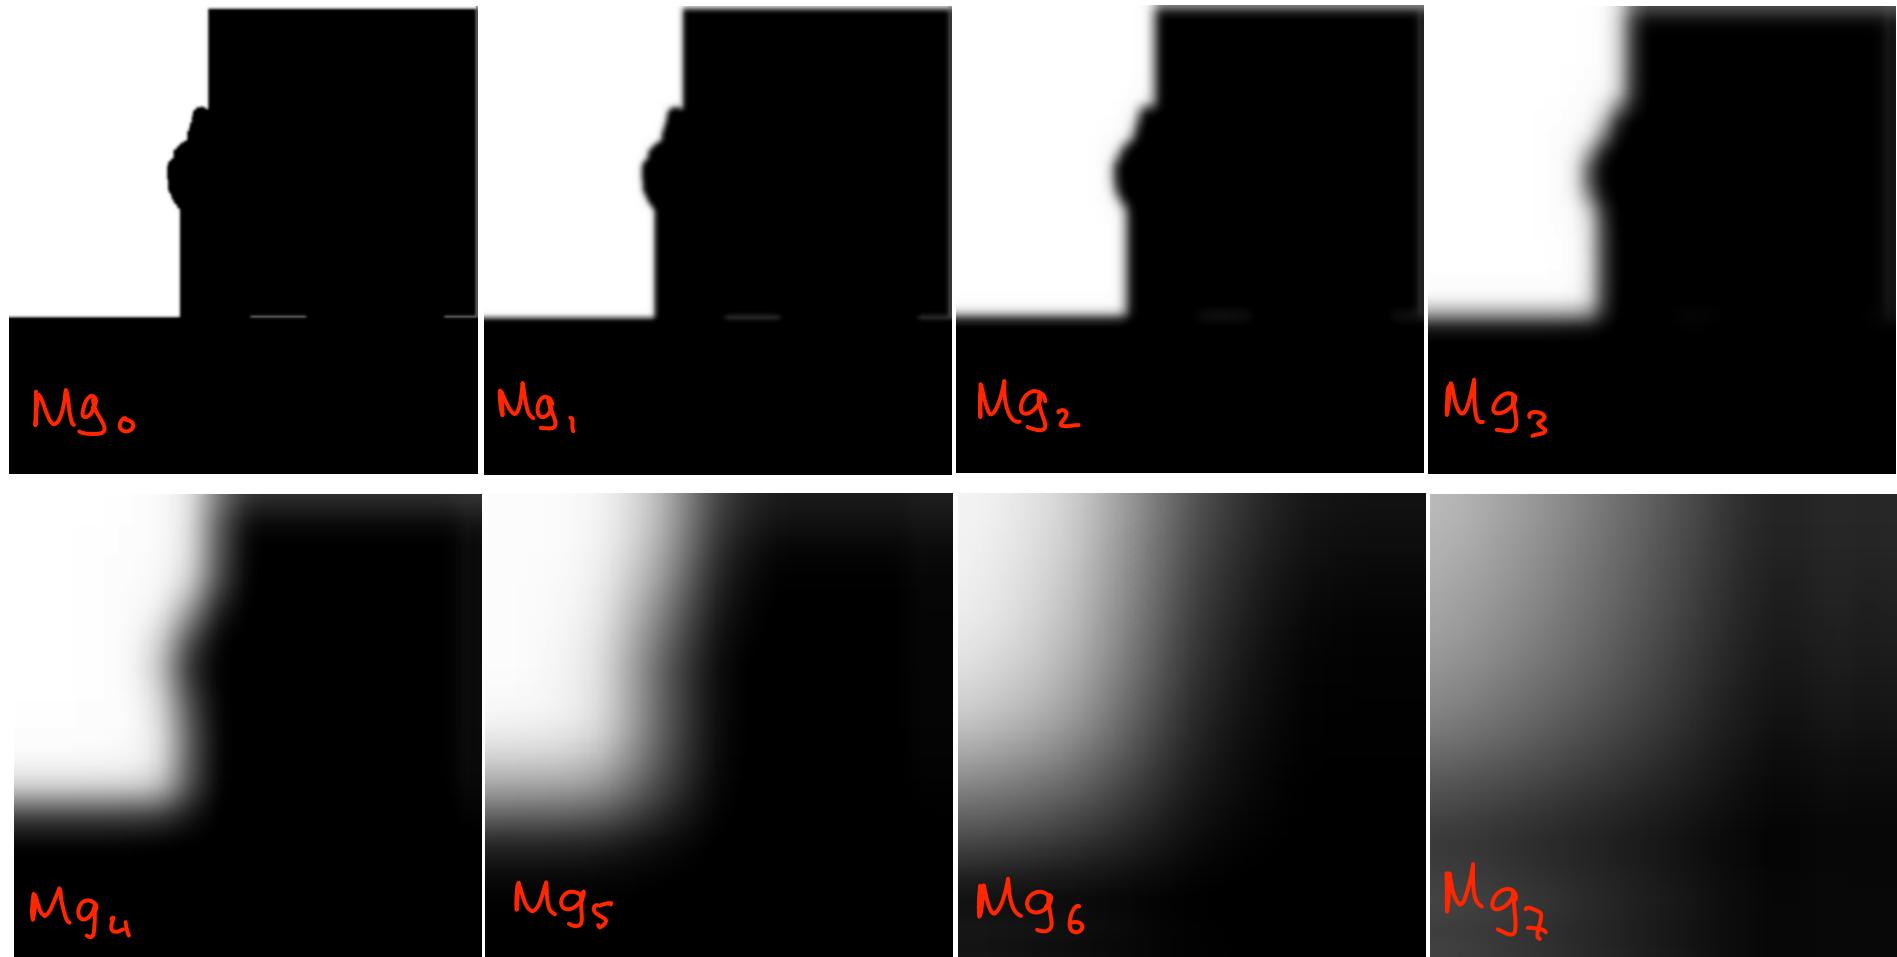


Pyramid Blending Algorithm

The algorithm effectively uses a different alpha matte for each level of detail

\uparrow detail \Rightarrow \downarrow feathering window

\downarrow detail \Rightarrow \uparrow feathering window



Pyramid Blending Algorithm

⑤ Compute level 0 of the Gaussian pyramid of S from $S_{L_0}, \dots, S_{L_{N-1}}, S_{g_N}$

Result S'



Pyramid Blending Algorithm



Blending Mis-Matched Photos (still looks OK)



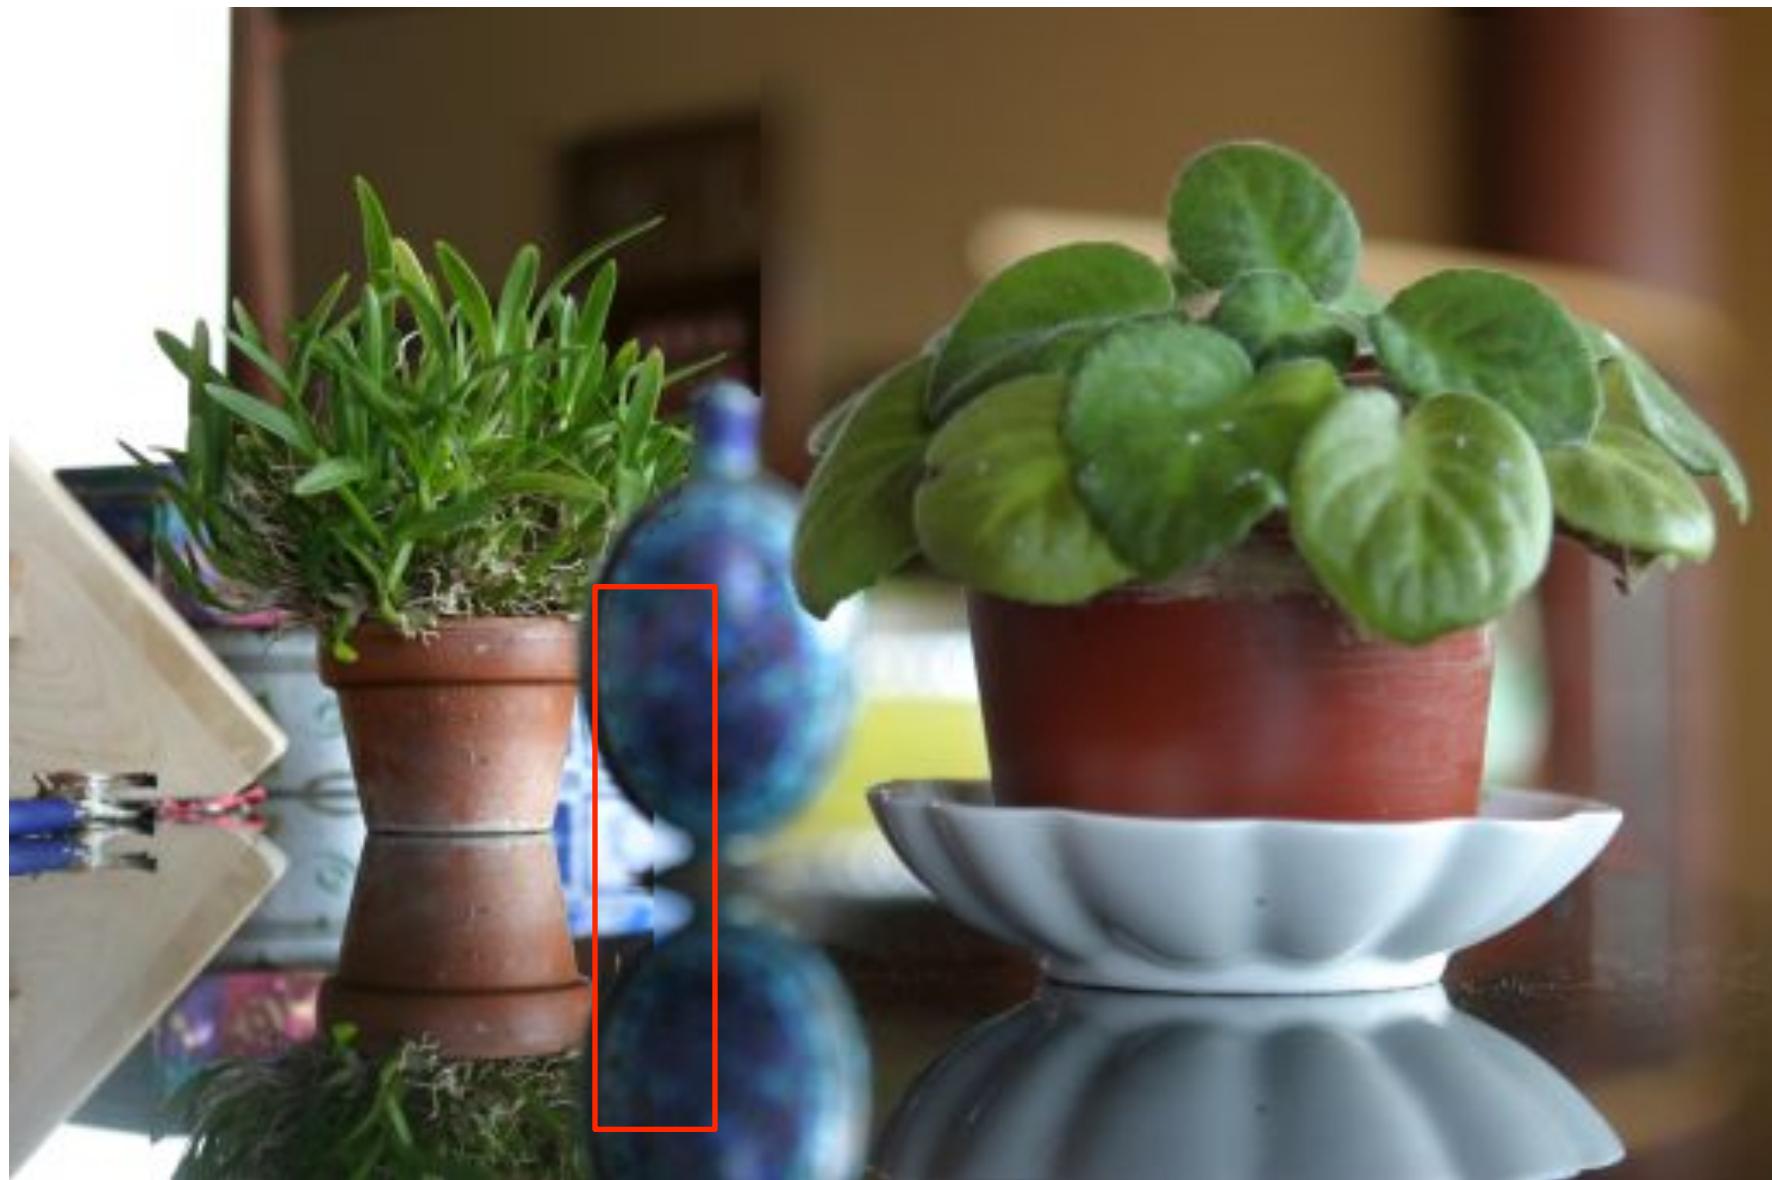
Merging Mis-Matched Photos (no blend)



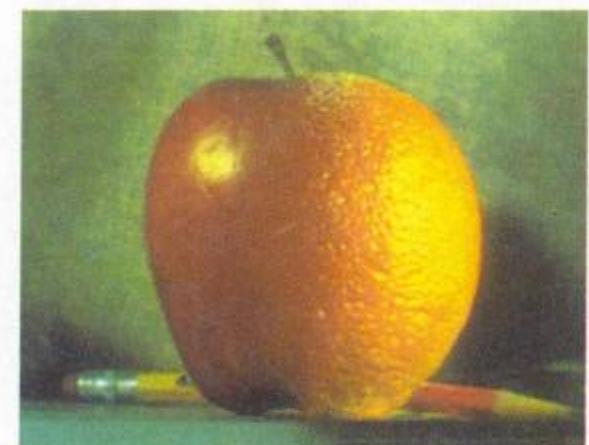
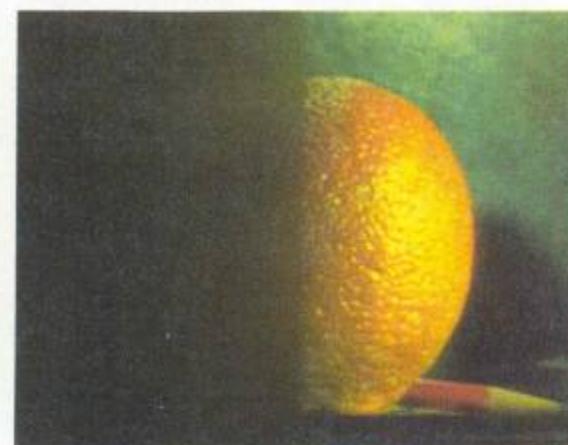
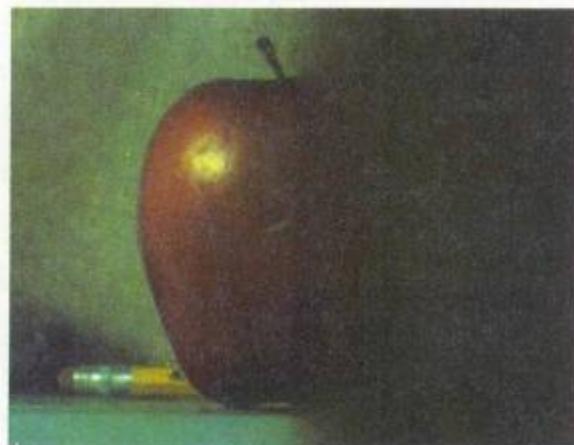
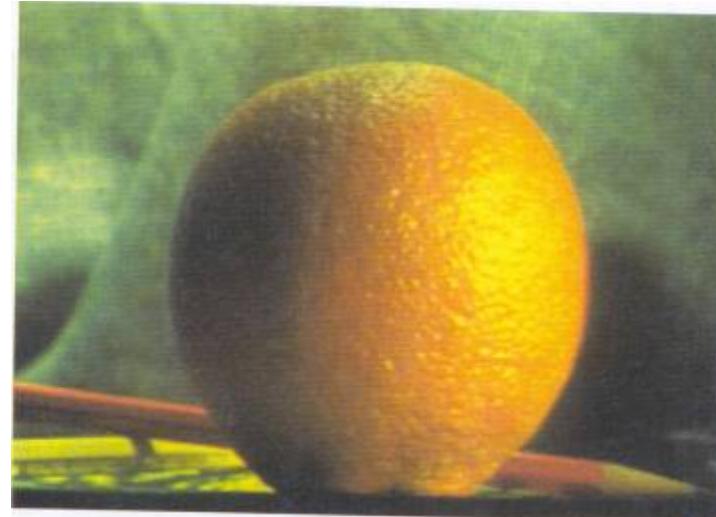
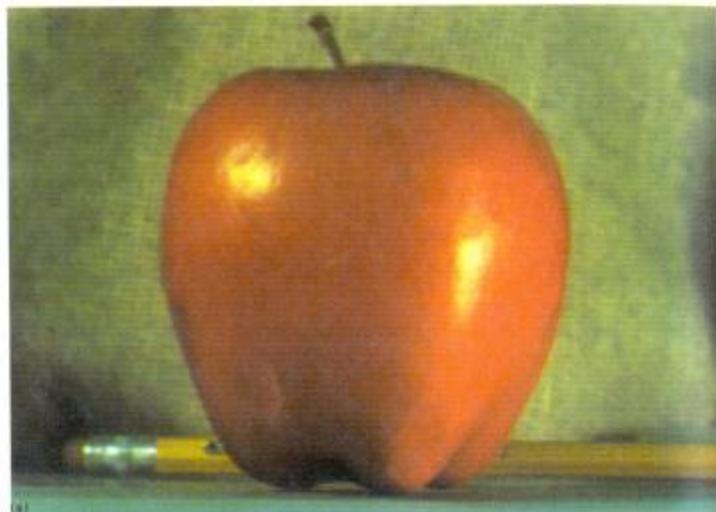
Blending without Using Pyramid



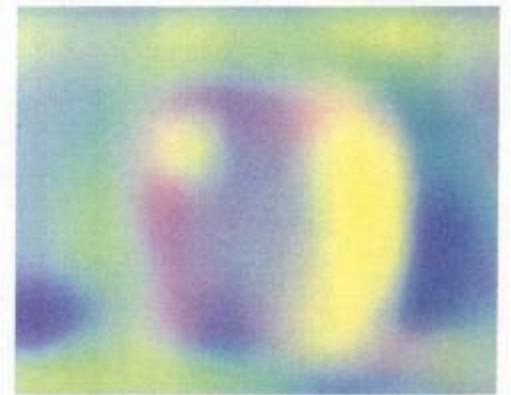
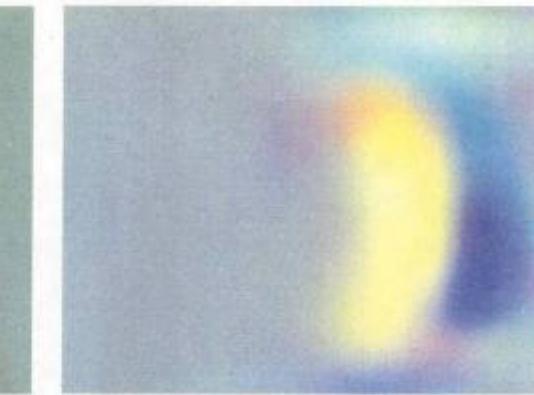
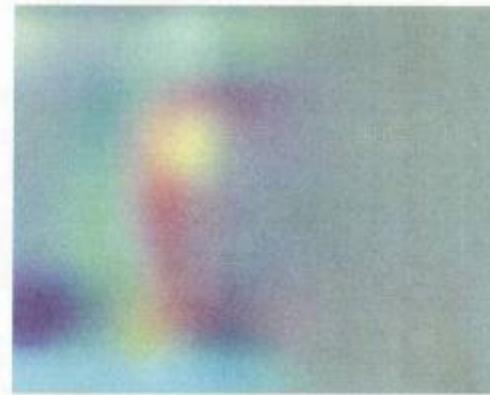
Blending without Using Pyramid



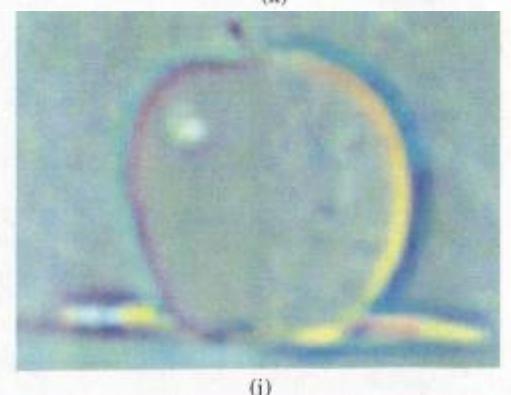
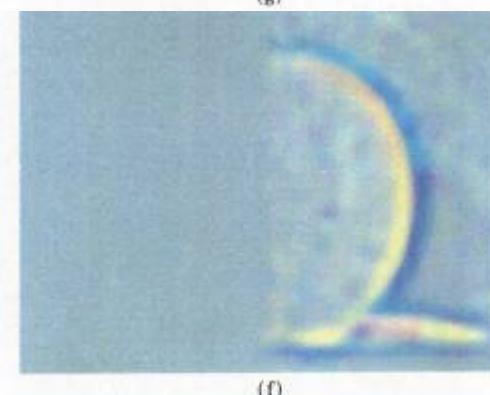
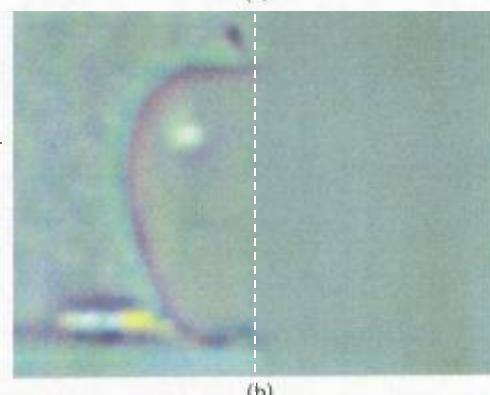
Pyramid Blending



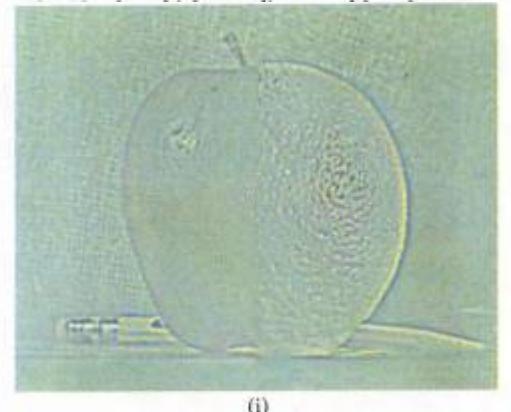
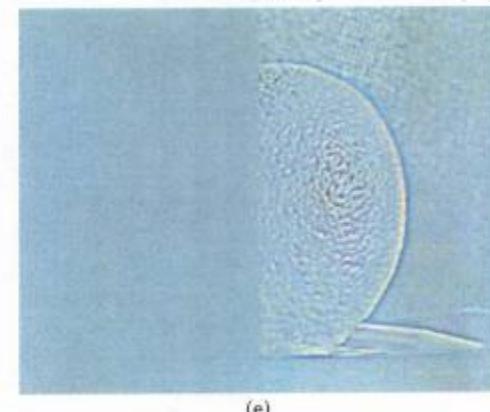
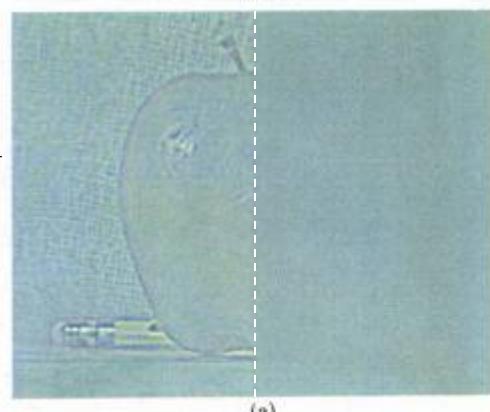
laplacian
level
4



laplacian
level
2



laplacian
level
0

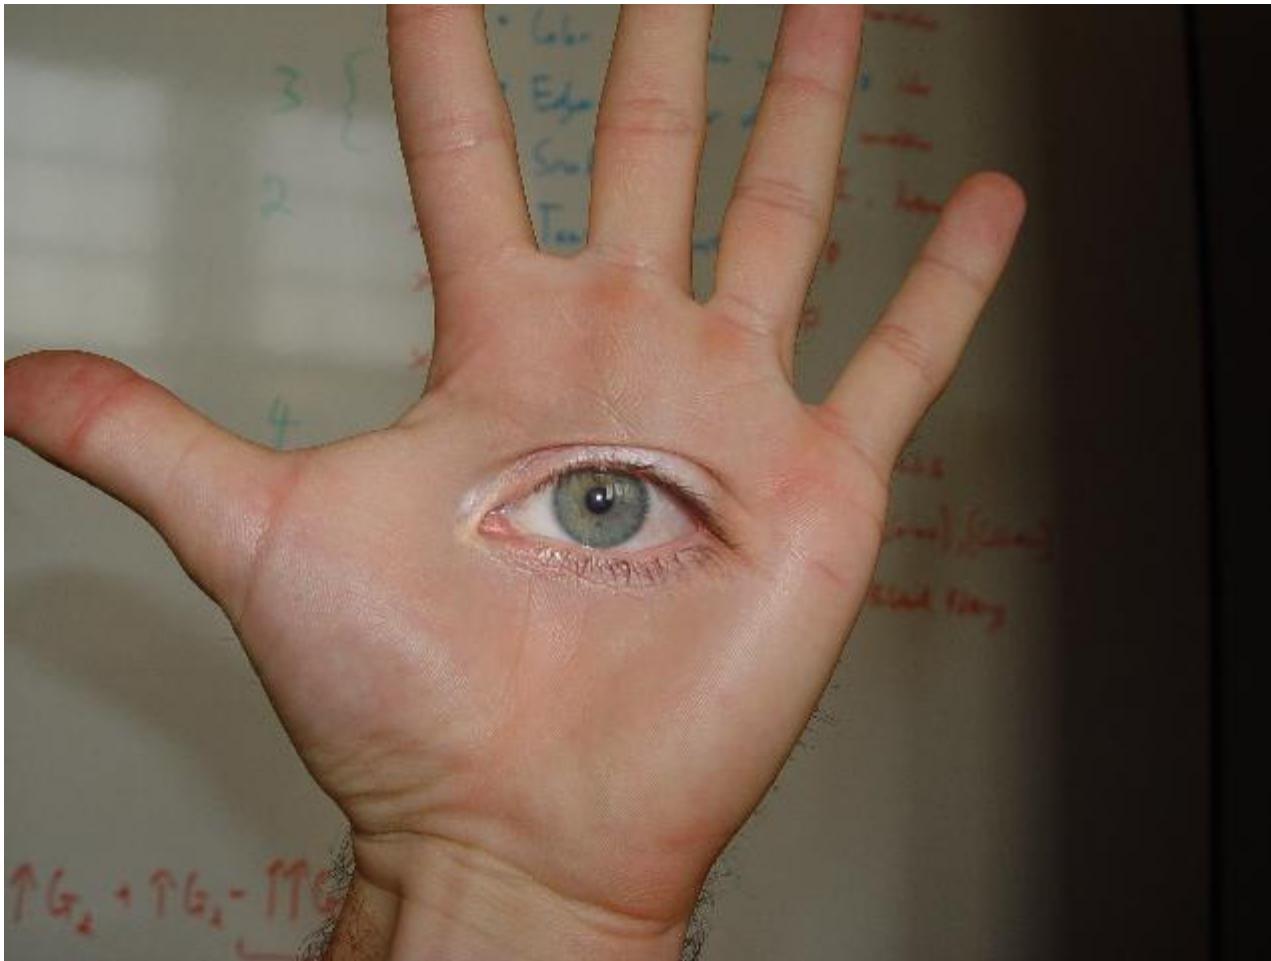


left pyramid

right pyramid

blended pyramid

Horror Photo



© prof. dmartin